

UNIVERSITÀ DEGLI
STUDI DI PADOVA

A Comparative Study of Pre-trained Deep Learning Models for Grapevine Leaf Disease Detection

Project presentation
(Vision and cognitive systems 2023-2024)

February 2024

Presented by

**Alessandro Canel
Lorenzo D'Antoni
Fateme Baghaei Saryazdi**

Abstract

- **Grape leaf disease classification with deep-learning approach**
- 4 classes: **Black Rot, ESCA, LeafBlight, Healthy**
- Enhancing a **Kaggle** project by using the **TensorFlow** framework and **adapted** it to the **PyTorch** library
- **evaluated their performance**



Introduction

Motivation:

- **Grapes** are one of the **most cultivated** fruits, and they have **many diseases**
- **Early detection** is crucial for the wine industry.

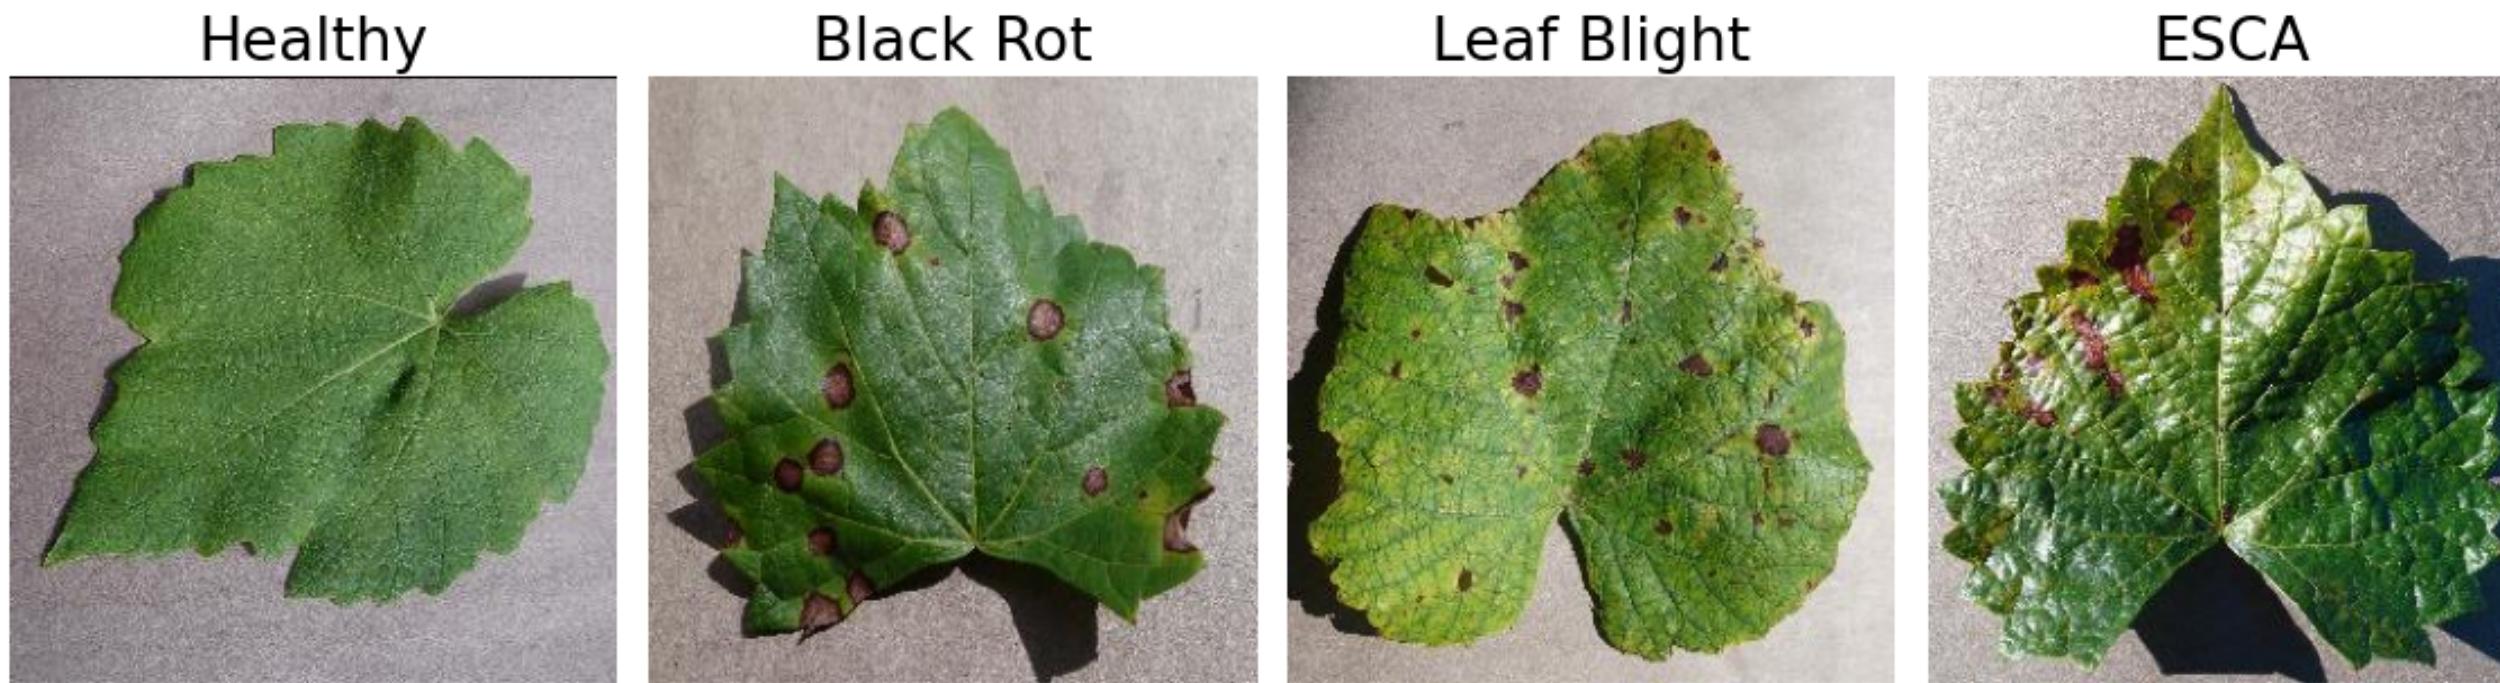
Challenge:

- Human inspection is **time-consuming**.

Solution:

- **Computer vision** and **machine learning** approaches can automate and **improve** the process

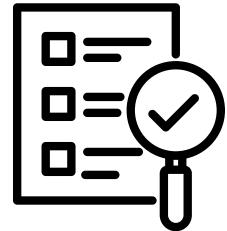




Grape leaf disease classification using deep learning, identifying whether the leaf is **infected** by **Black Rot, ESCA, Leaf Blight** or is **healthy**.

Problems:

- **Scarcity** of annotated **datasets**
- Need for **real-time detection**
- **Diversity** and **Complexity** of **disease** spots



Results

Achieved high accuracy on the test set, but struggled to distinguish between **Black Rot** and **ESCA**, which are similar diseases





Related works

Before 2016:

- **K-means** clustering:
 - find diseased regions (segmentation)
- extract **color** and **texture** features
- **PCA**
 - reduce the dimensionality of the plant image features and retain the most important ones
- **SVM**:
 - detect the type of leaf disease



Problems

- manually engineered features:
 - **Time-consuming**
 - **Incomplete**
 - **No generalization**
- **extensive preprocessing** to deal with image variations and **noise**

After 2016: CNN

- automatically learn **discriminative** features from **images**
- avoiding **complicated** preprocessing
- **preserving** relevant **information**
- capture **hierarchical** spatial **features**
- extracting **complex patterns** and structures
- computationally **efficient** (parameter sharing)
- **translation invariant**

Two Main Papers



Faster R-CNN and ResNet:

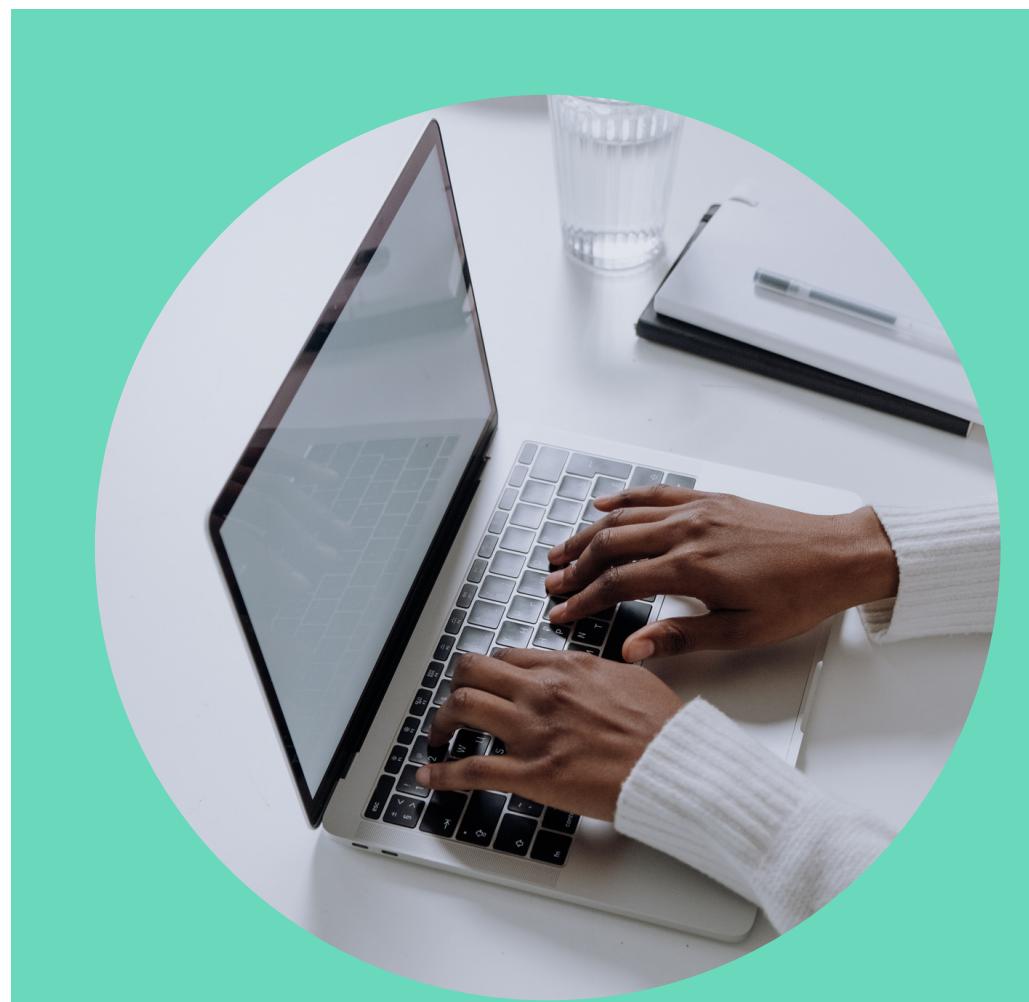
- enhanced the base model with **state-of-the-art** deep-learning modules
- **improving** the multiscale diseased **spot extraction** and **detection speed**
- **trained** and tested the **model** with **simple** and **complex backgrounds**
- **augmented** the **dataset** using **digital image processing** (to prevent the overfitting problem)

YOLOXS model:

- **collected** and used a dataset of 15 leaf diseases in natural environments
- **improved** the **generalization** and **speed** of the existing models using an improved YOLOXS model

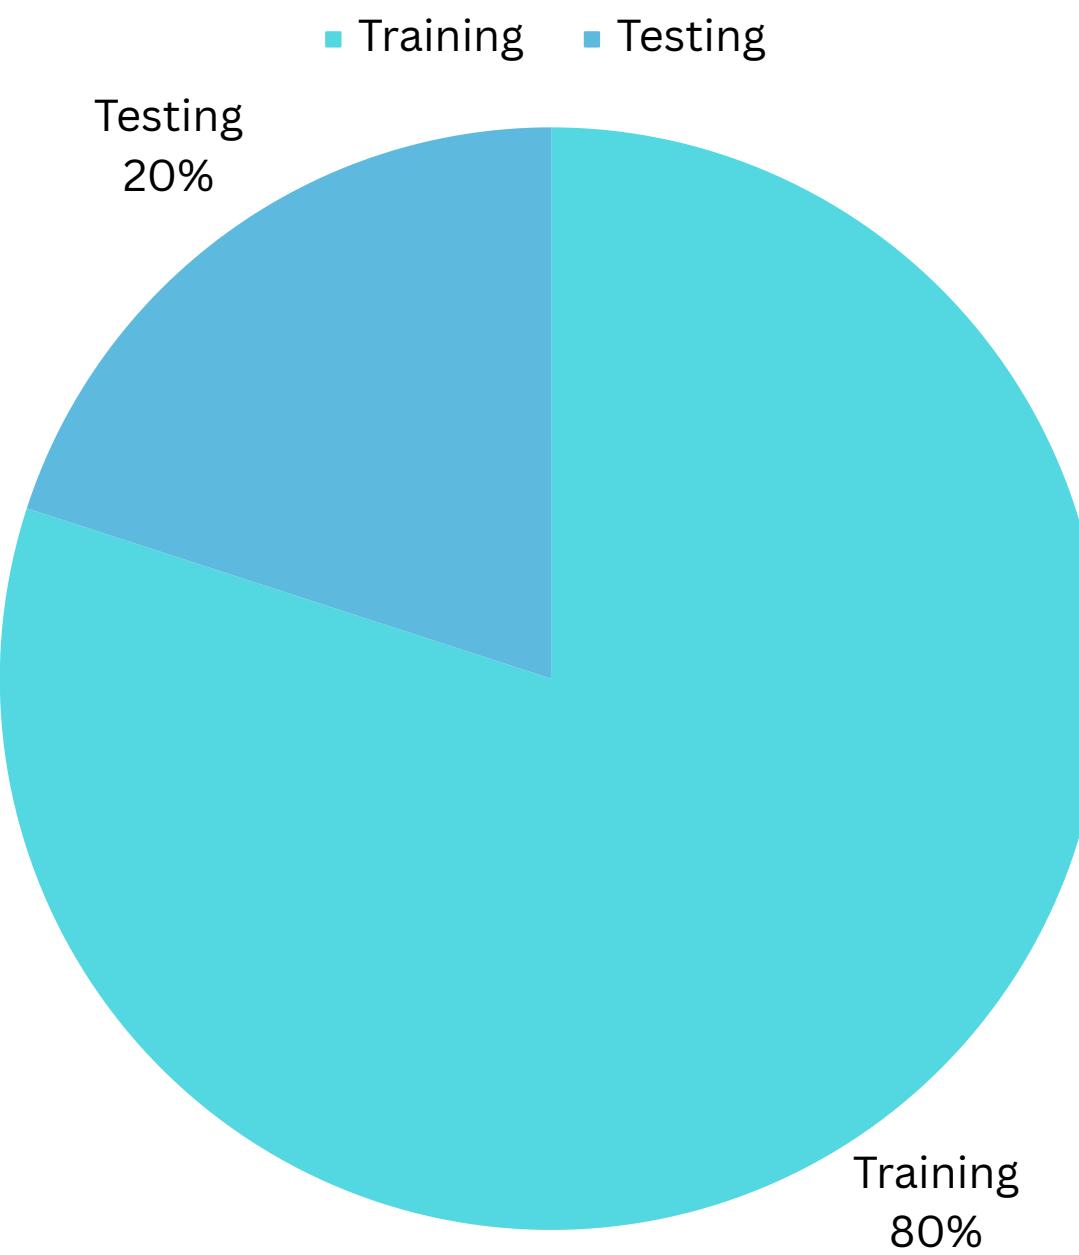
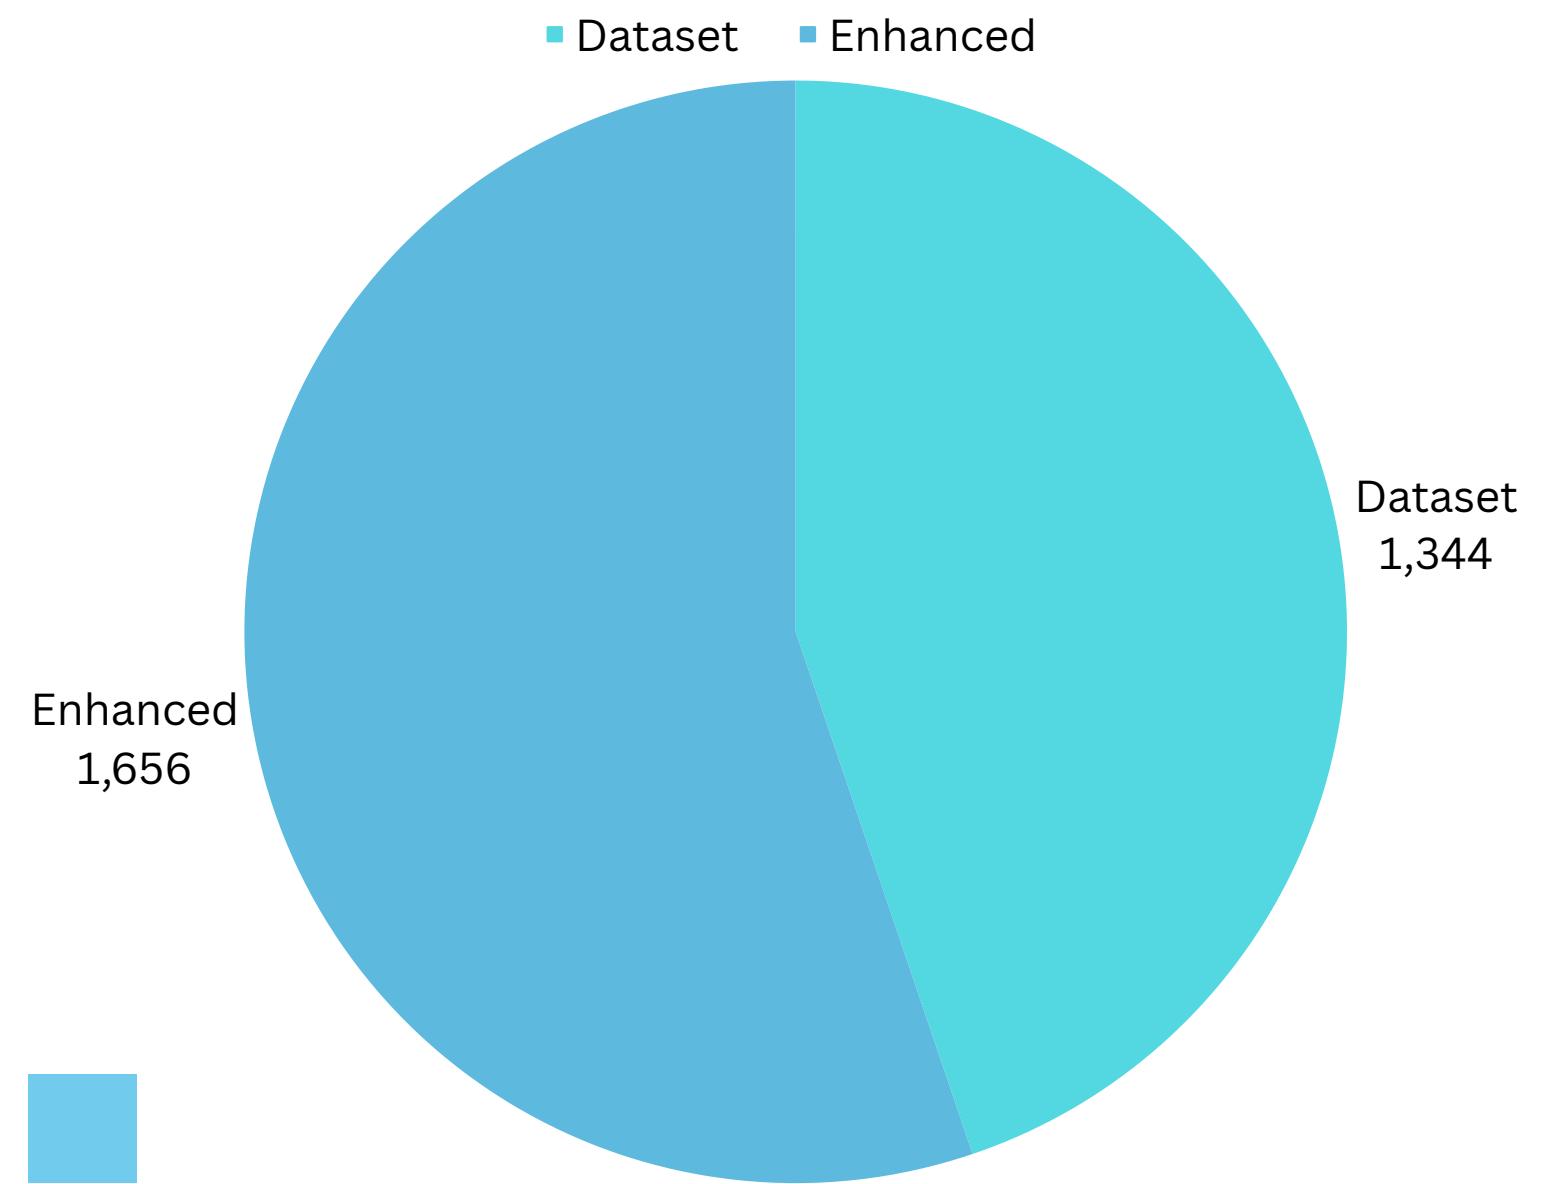
[Back to Agenda](#)

Dataset

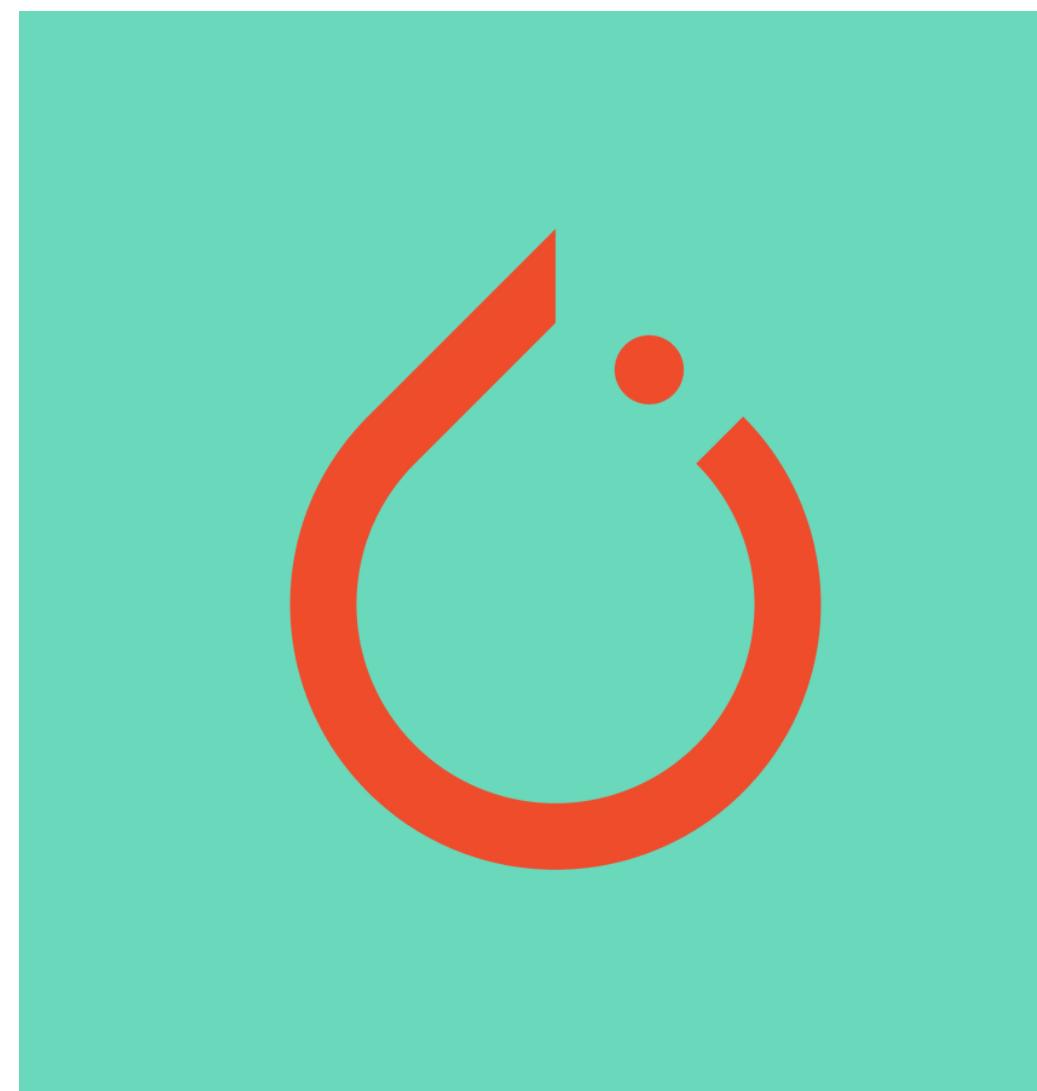


Enhancing

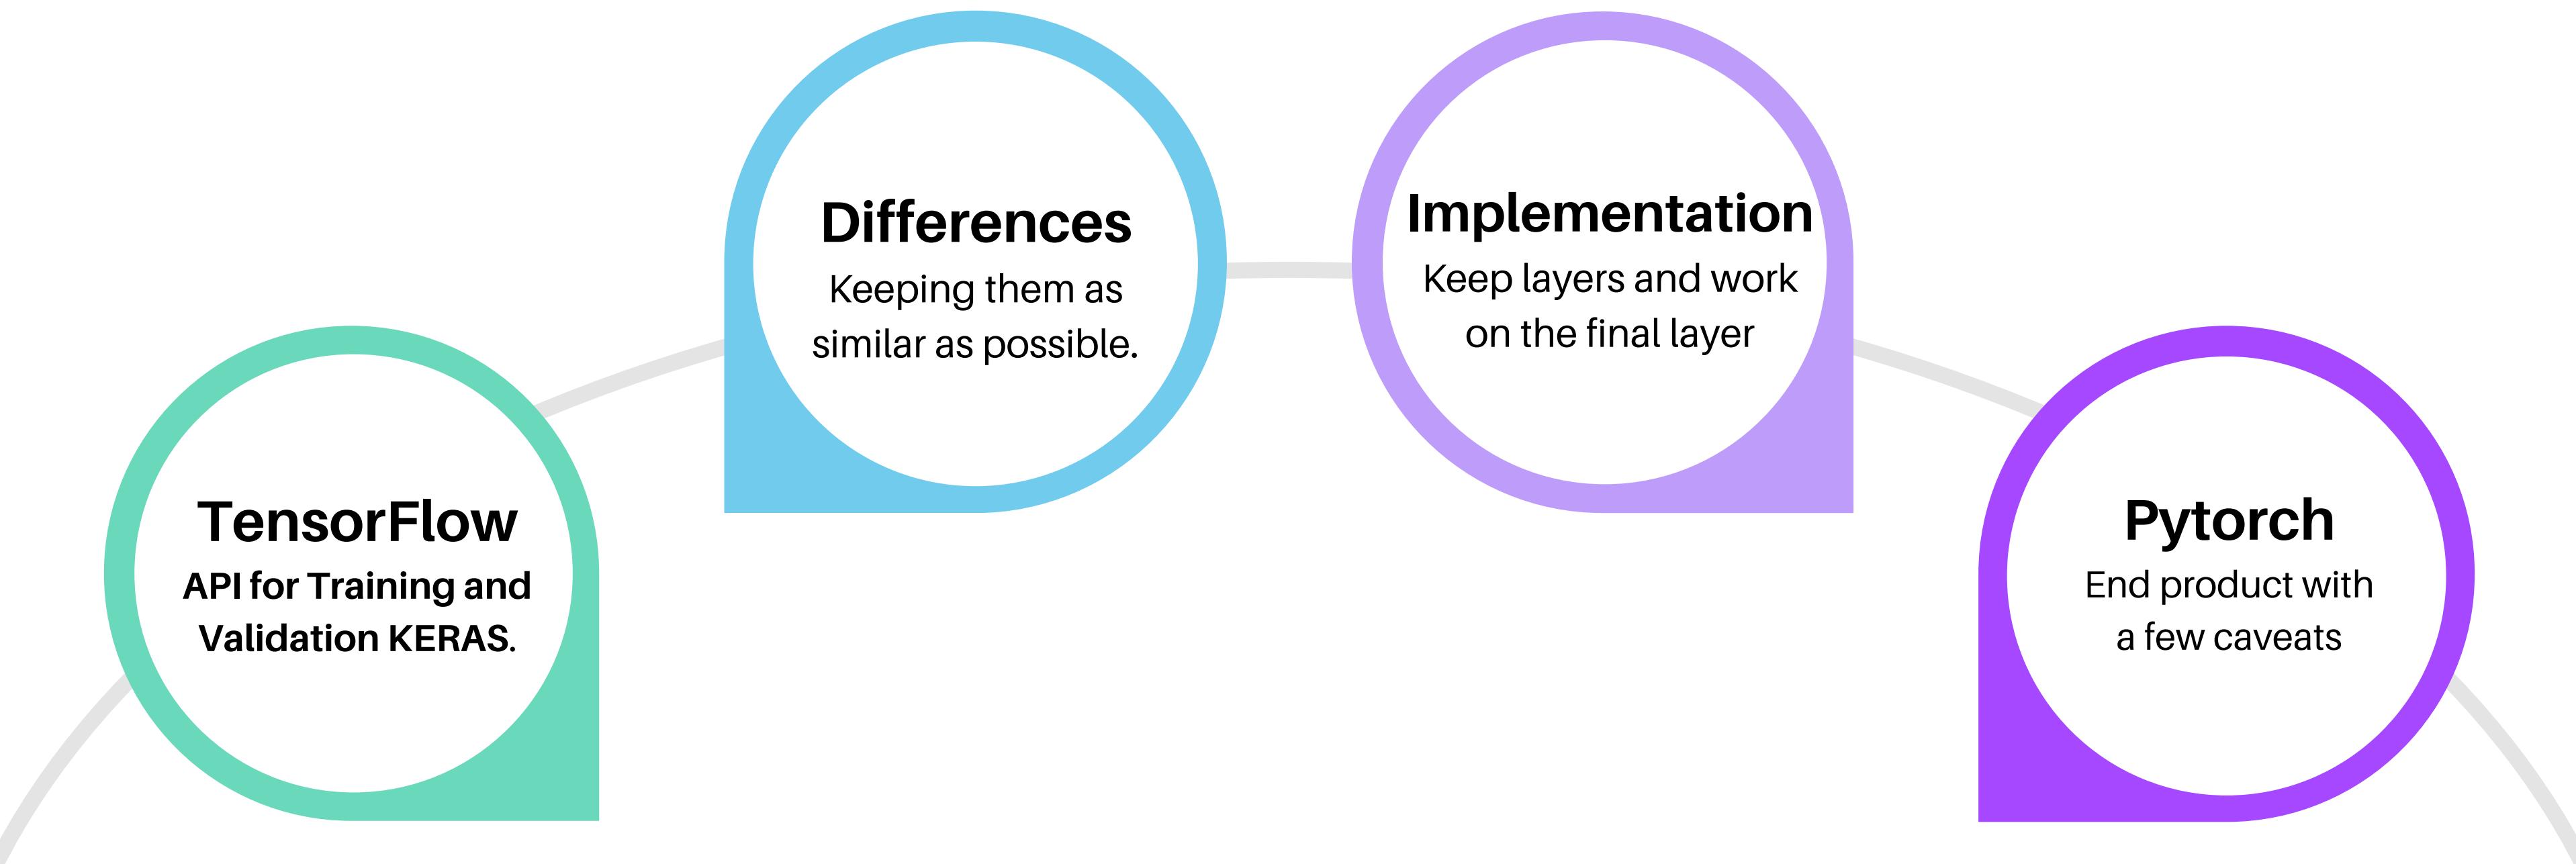
Understanding our Kaggle Dataset



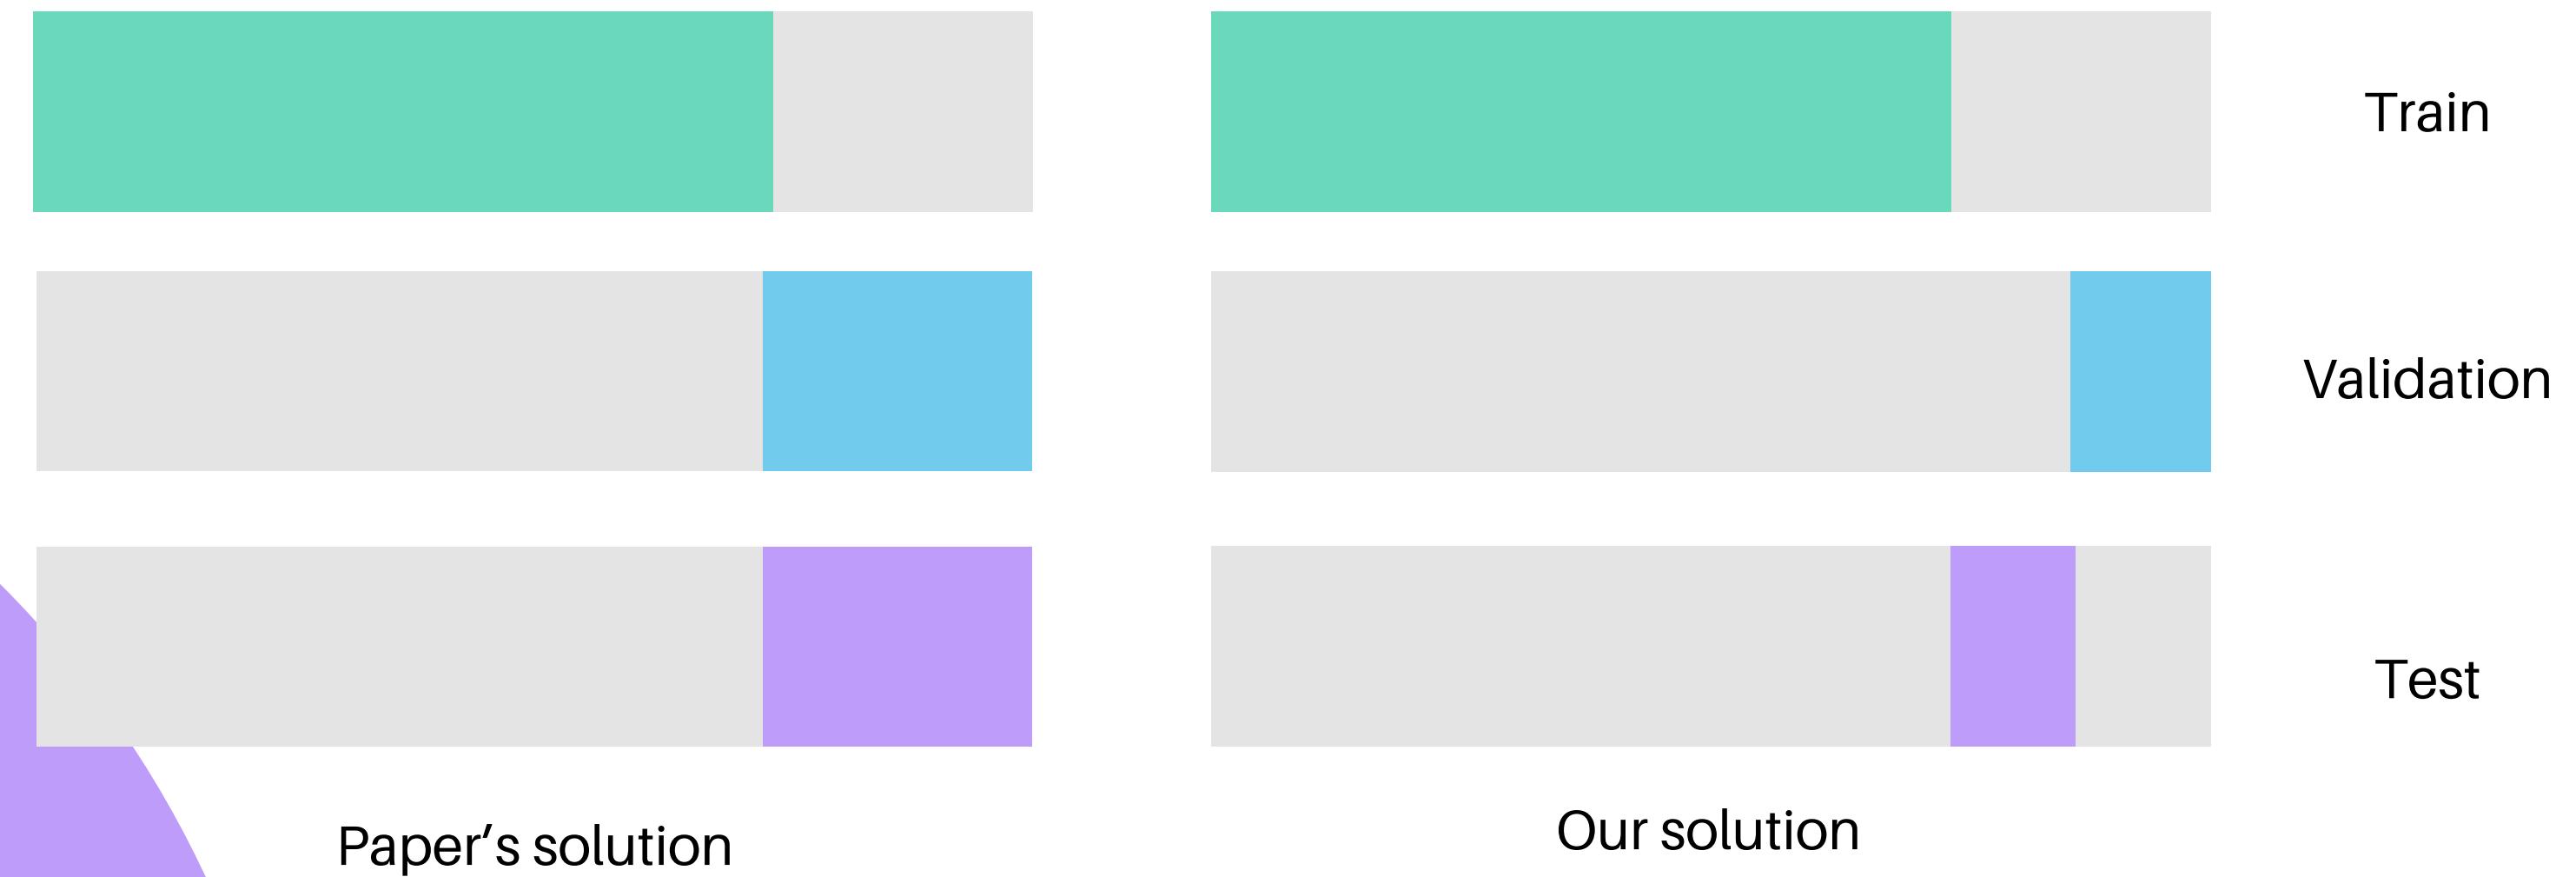
From TensorFlow to Pytorch



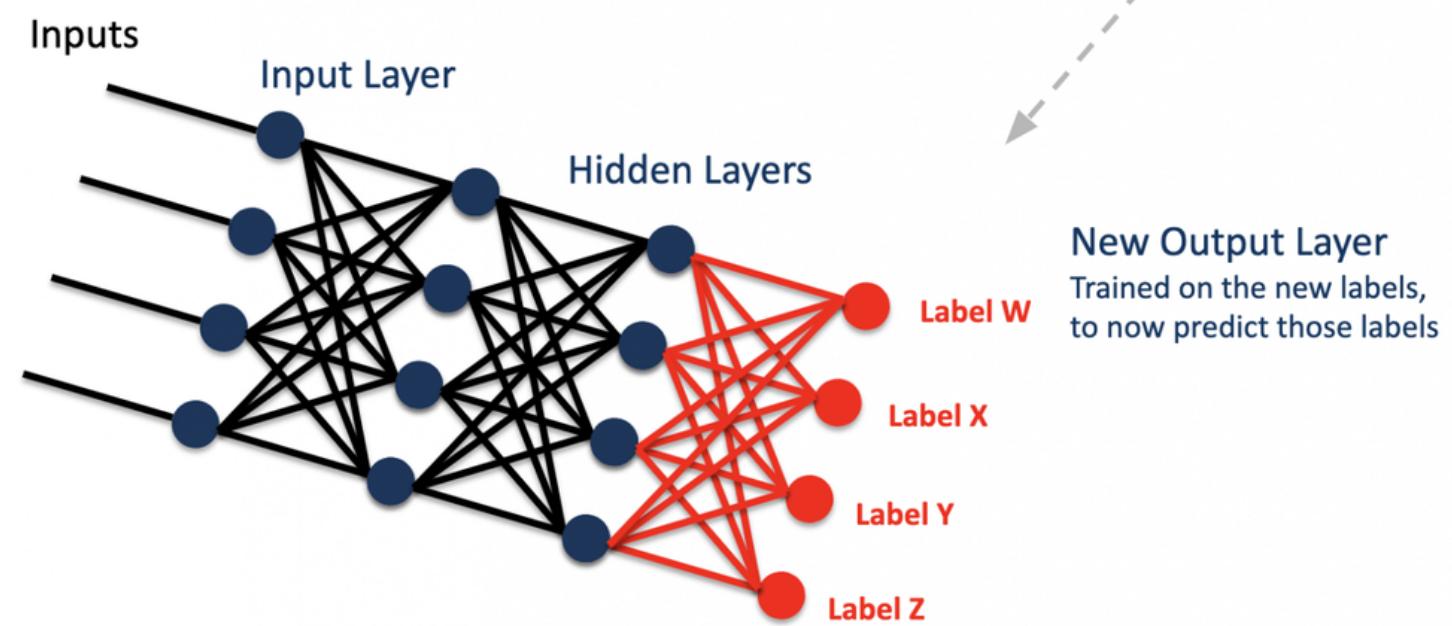
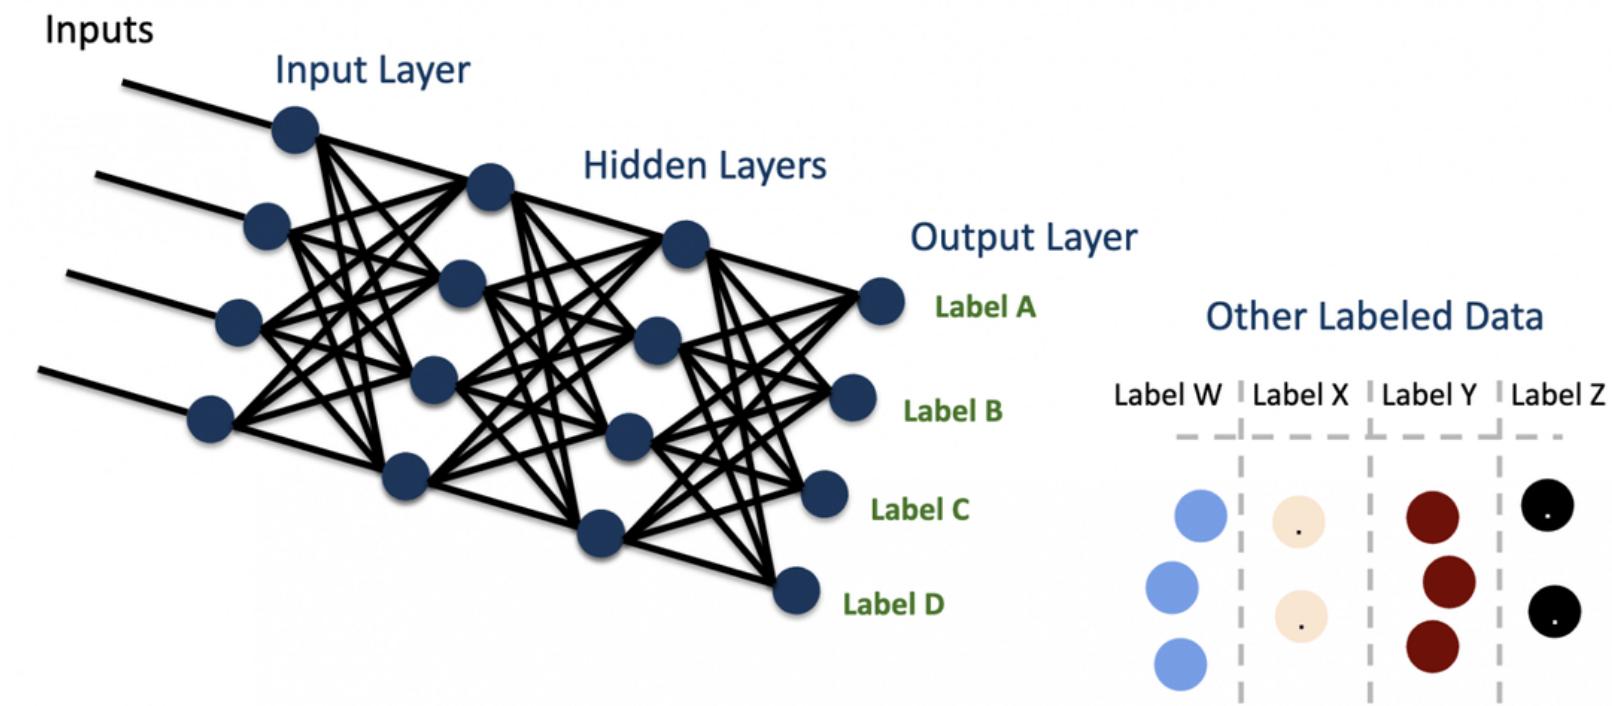
TensorFlow to Pytorch



Testing and Validation

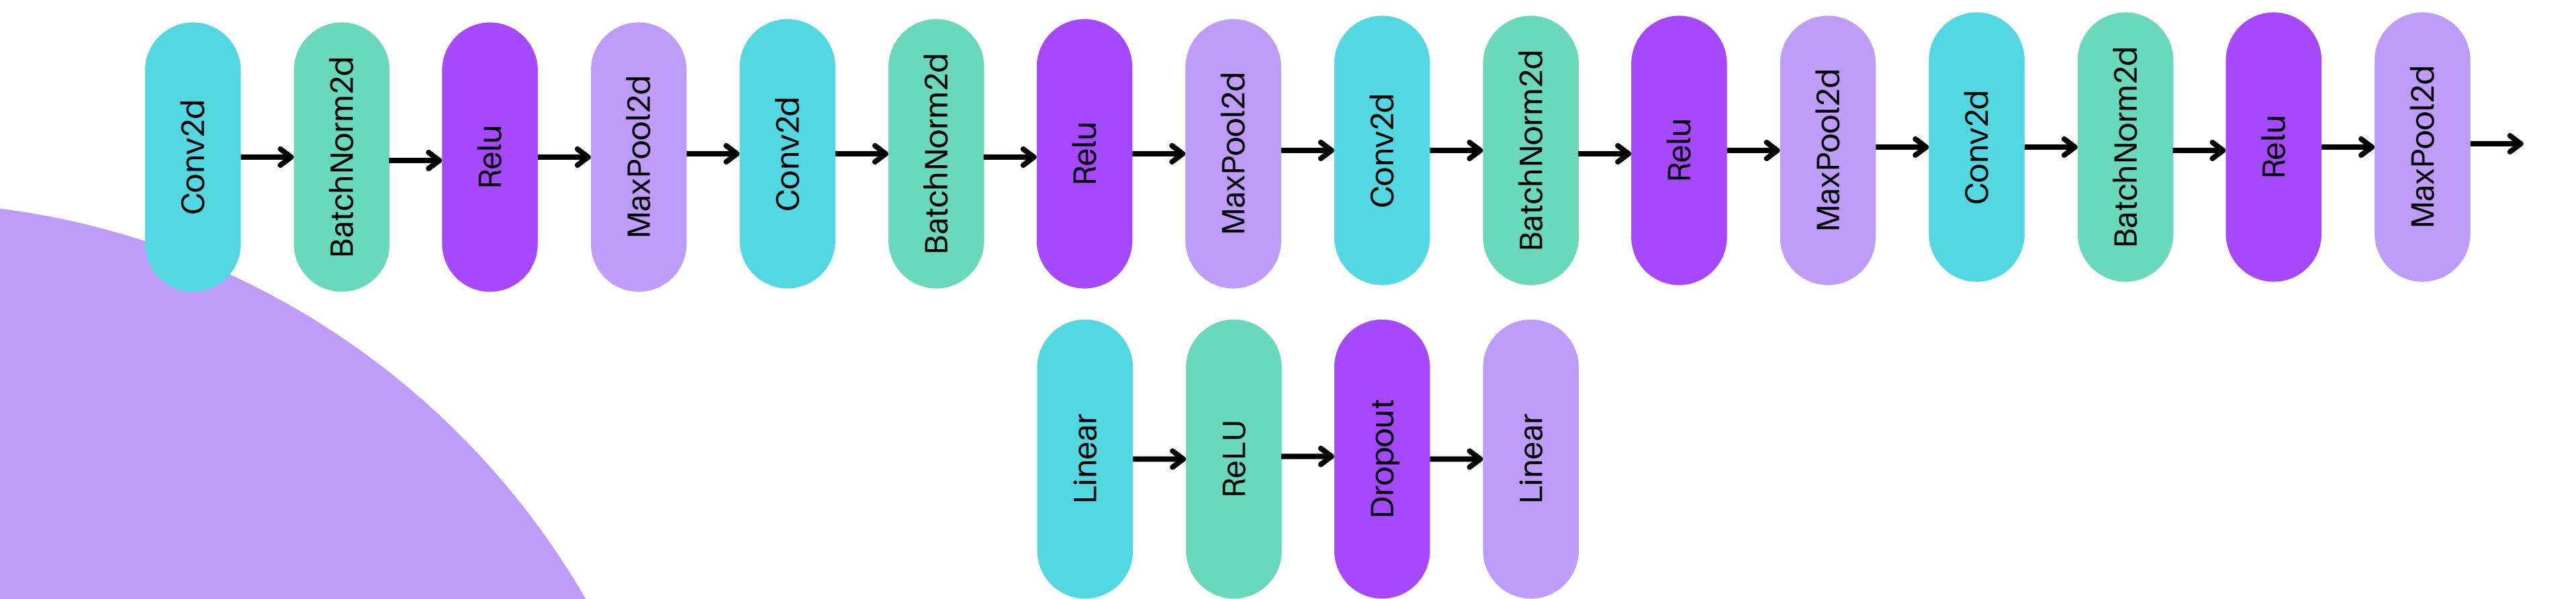


Transfer Learning

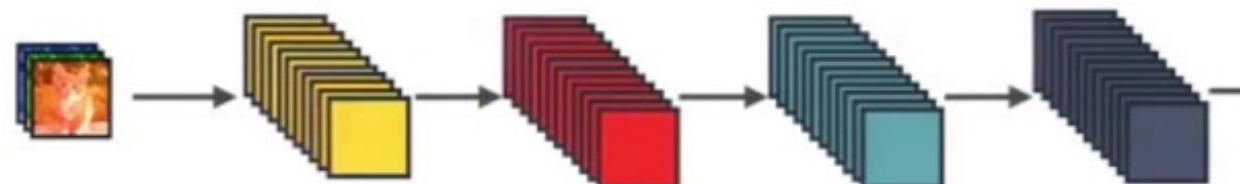


CNN	DenseNet	EfficientNet	MobileNet	ResNet	VGG	Ensemble model	
Baseline simple model	Layer-to-layer	Scaling along with parameters	Separable convolutions	Residual blocks	Deep CNN architecture		
	Transition Layers	Adapts to resource limits	Split convolution into depthwise and pointwise	Skip connections	Conv layers followed by max pooling		
	Control complexity and cost	Might overfit	Less intensive but good performance	Learns residual mappings	Reduces spatial dimensions		
	Many variations	Many variations	MobileNetV2, MobileNetV3, and MobileNetV3 Small	ResNet-18, ResNet-34, ResNet-50, ResNet 101 and ResNet-152	Authors used VGG16		
VERSION CHOSEN	Densenet-121	EfficientNet-B7	MobileNetV2	ResNet-50	VGG11	Best three models	

Our CNN

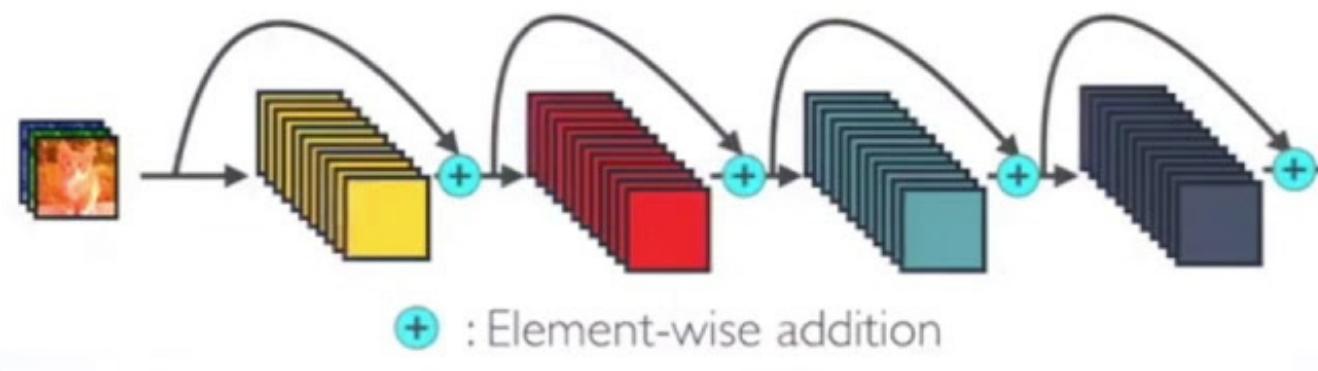


Standard Connectivity



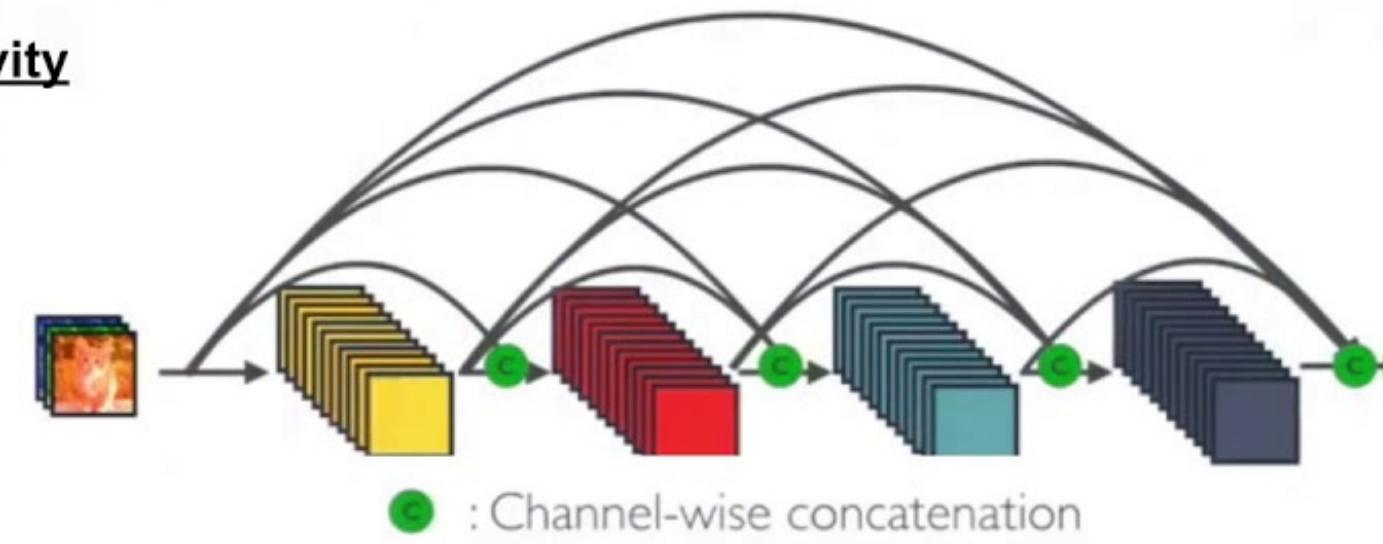
Successive convolutions

Resnet Connectivity



Element-wise feature summation
+ : Element-wise addition

DenseNet Connectivity



Feature concatenation
c : Channel-wise concatenation

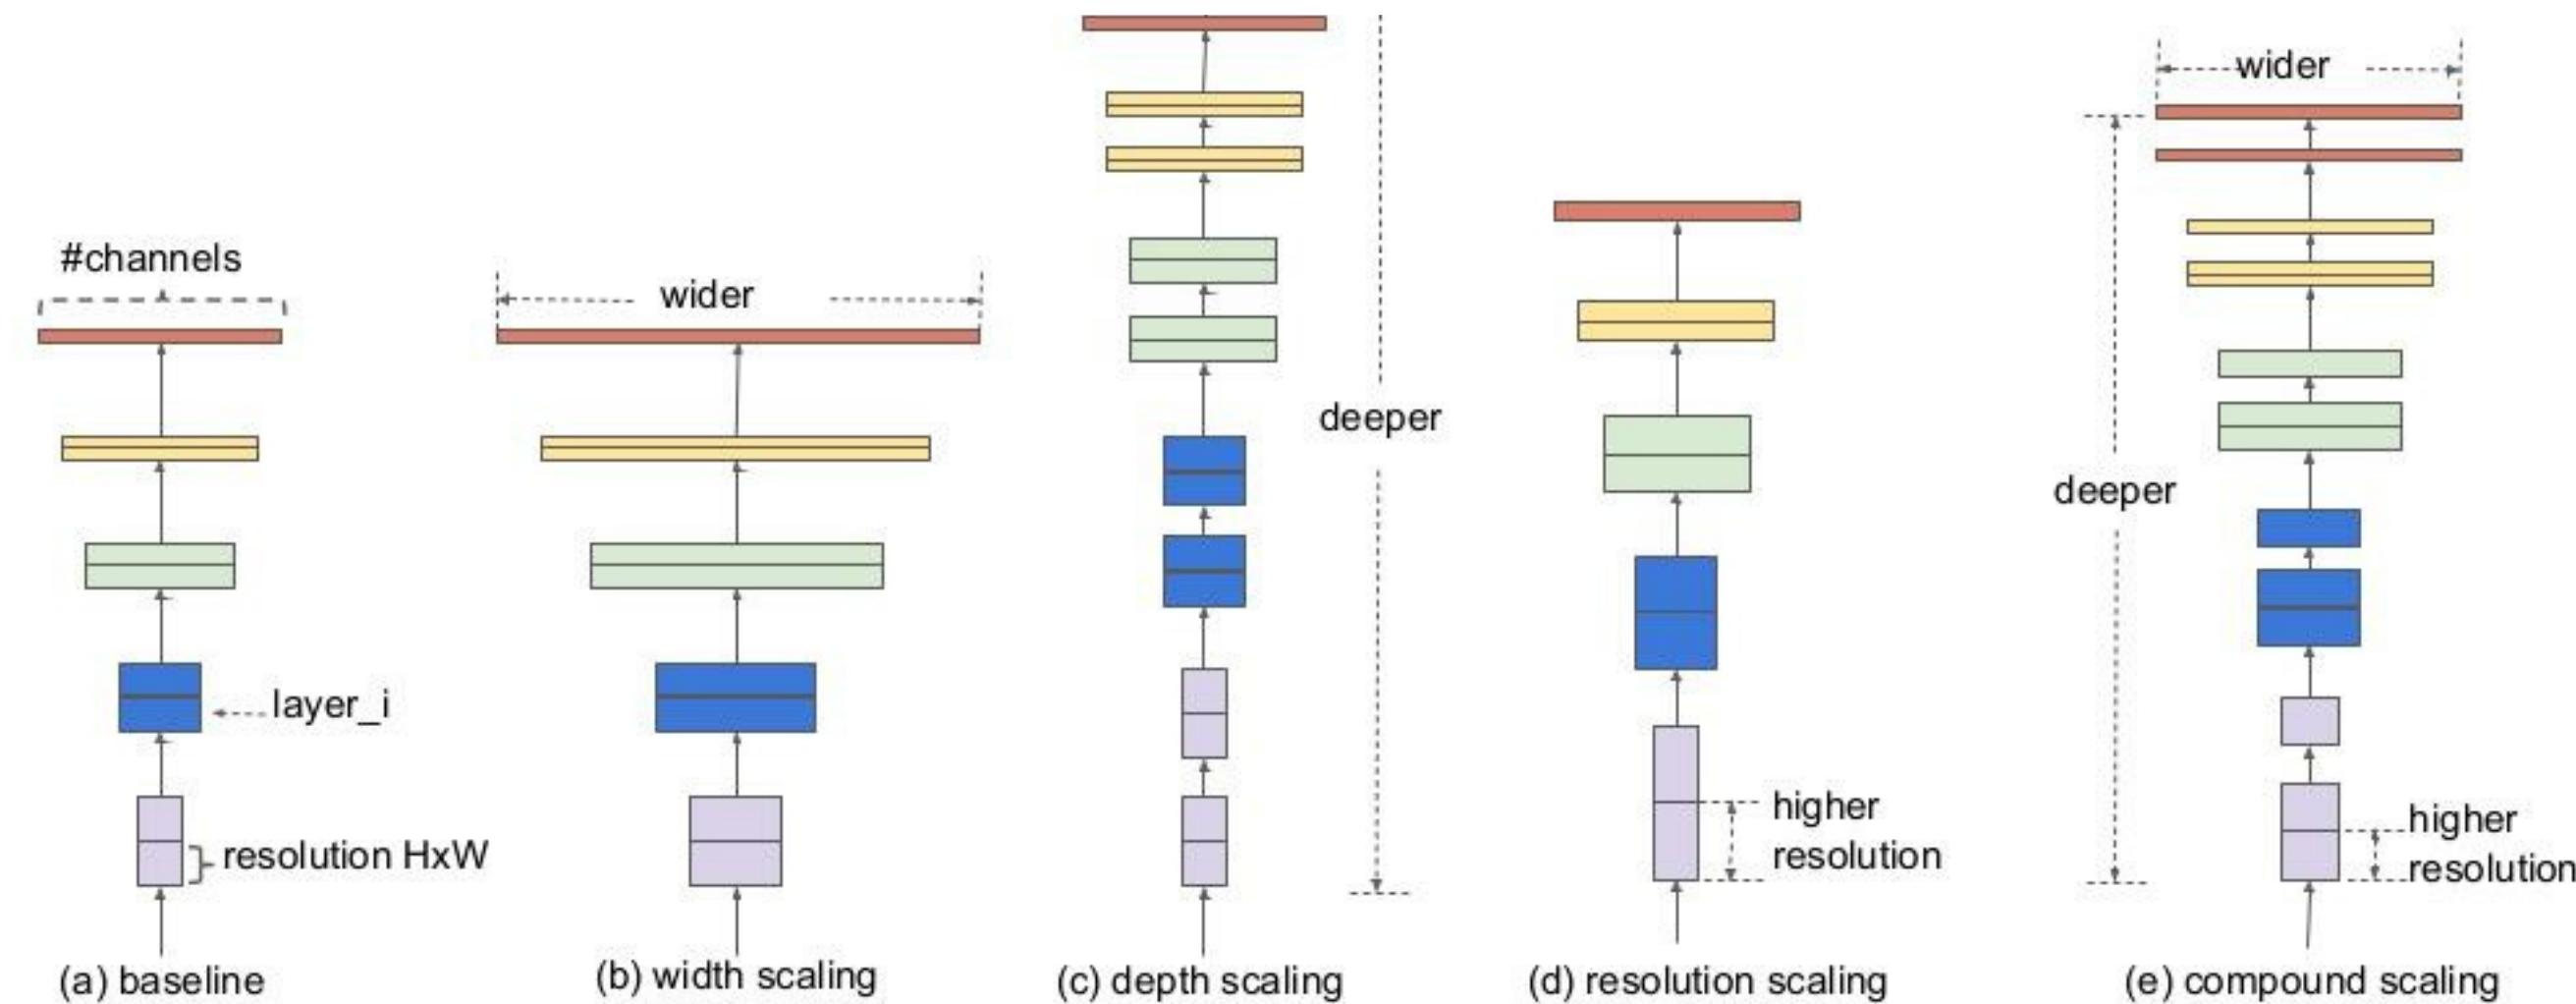
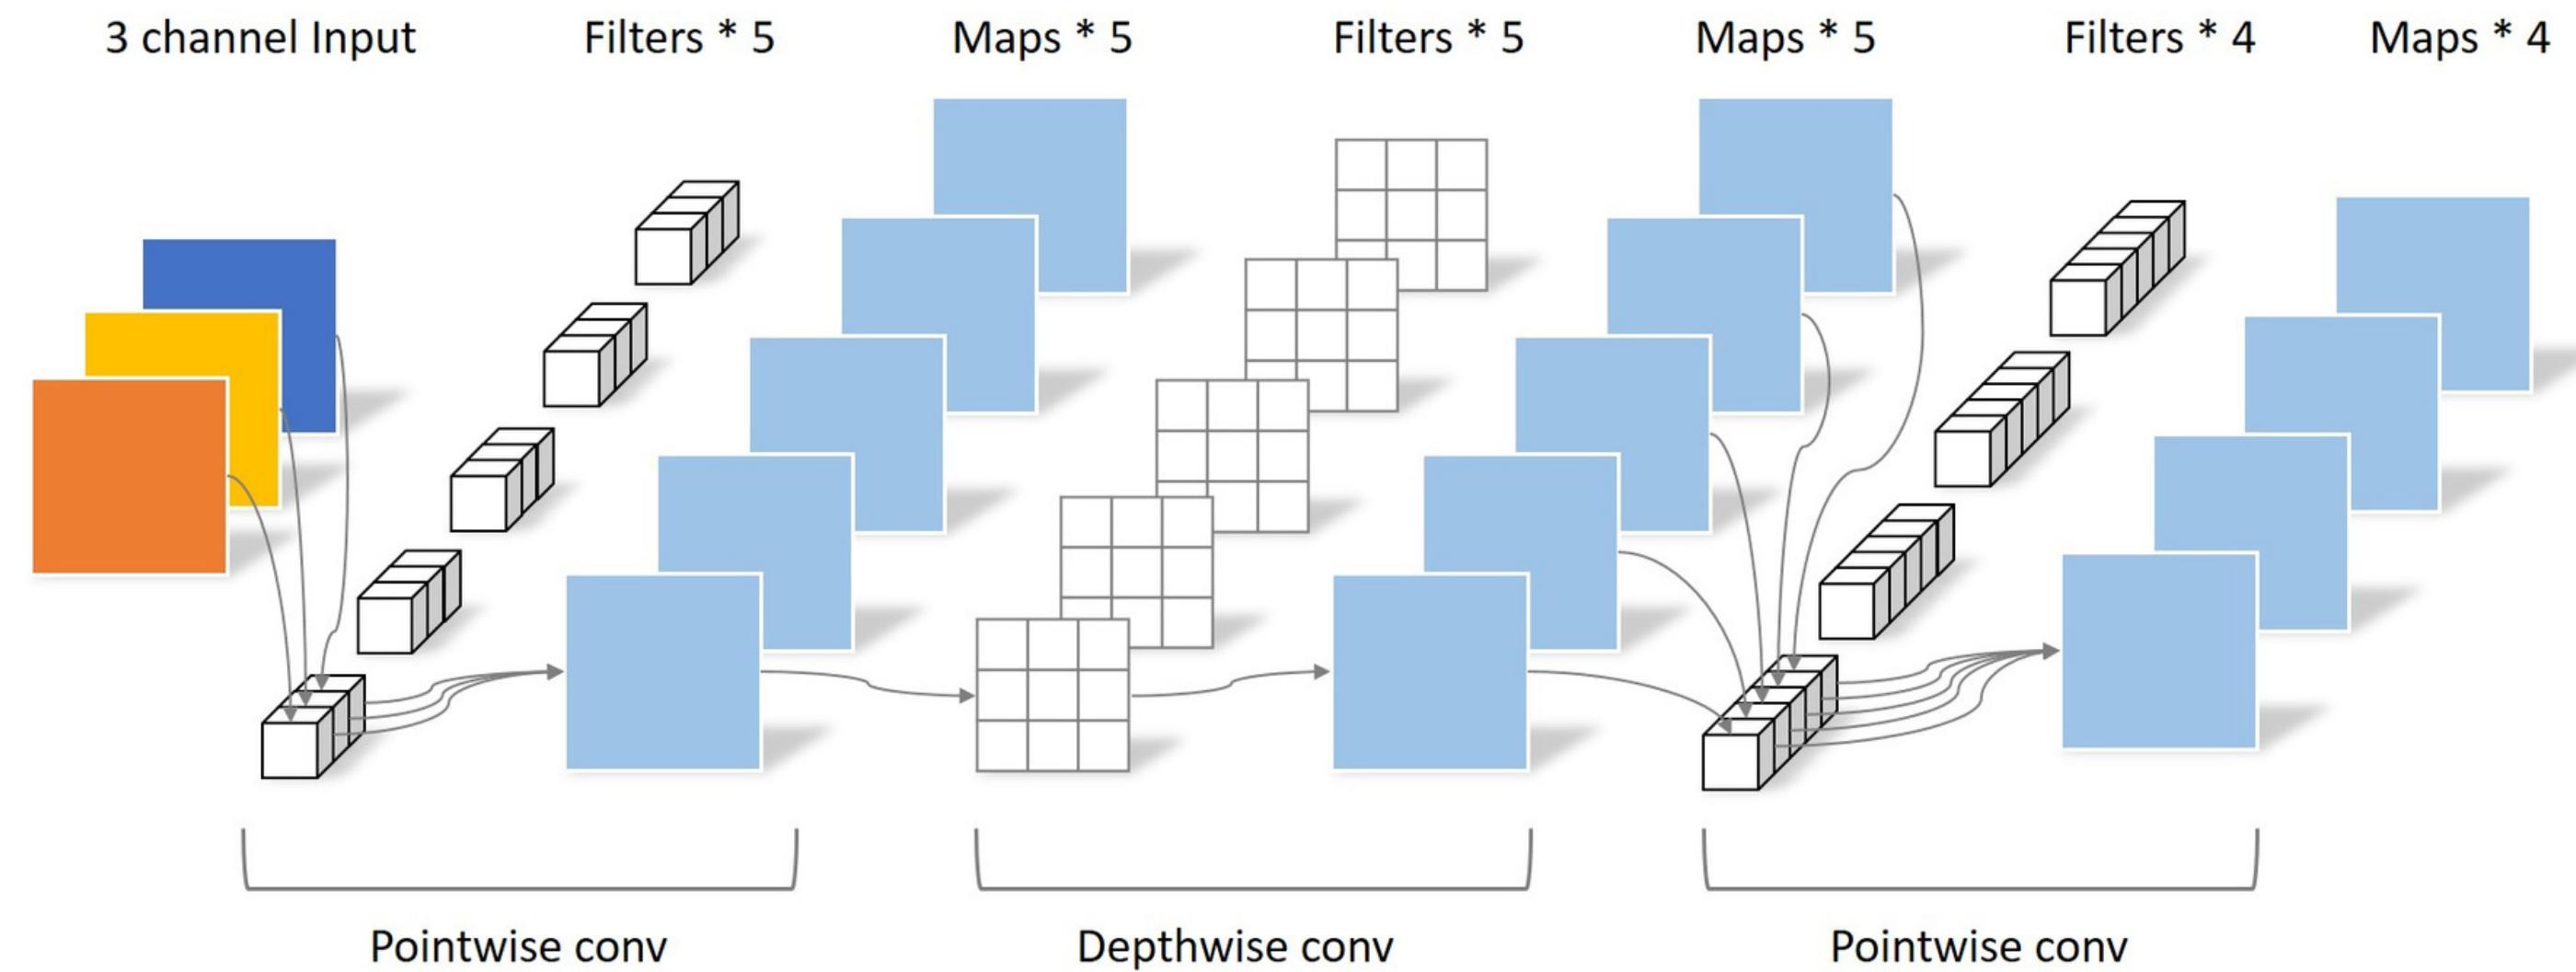
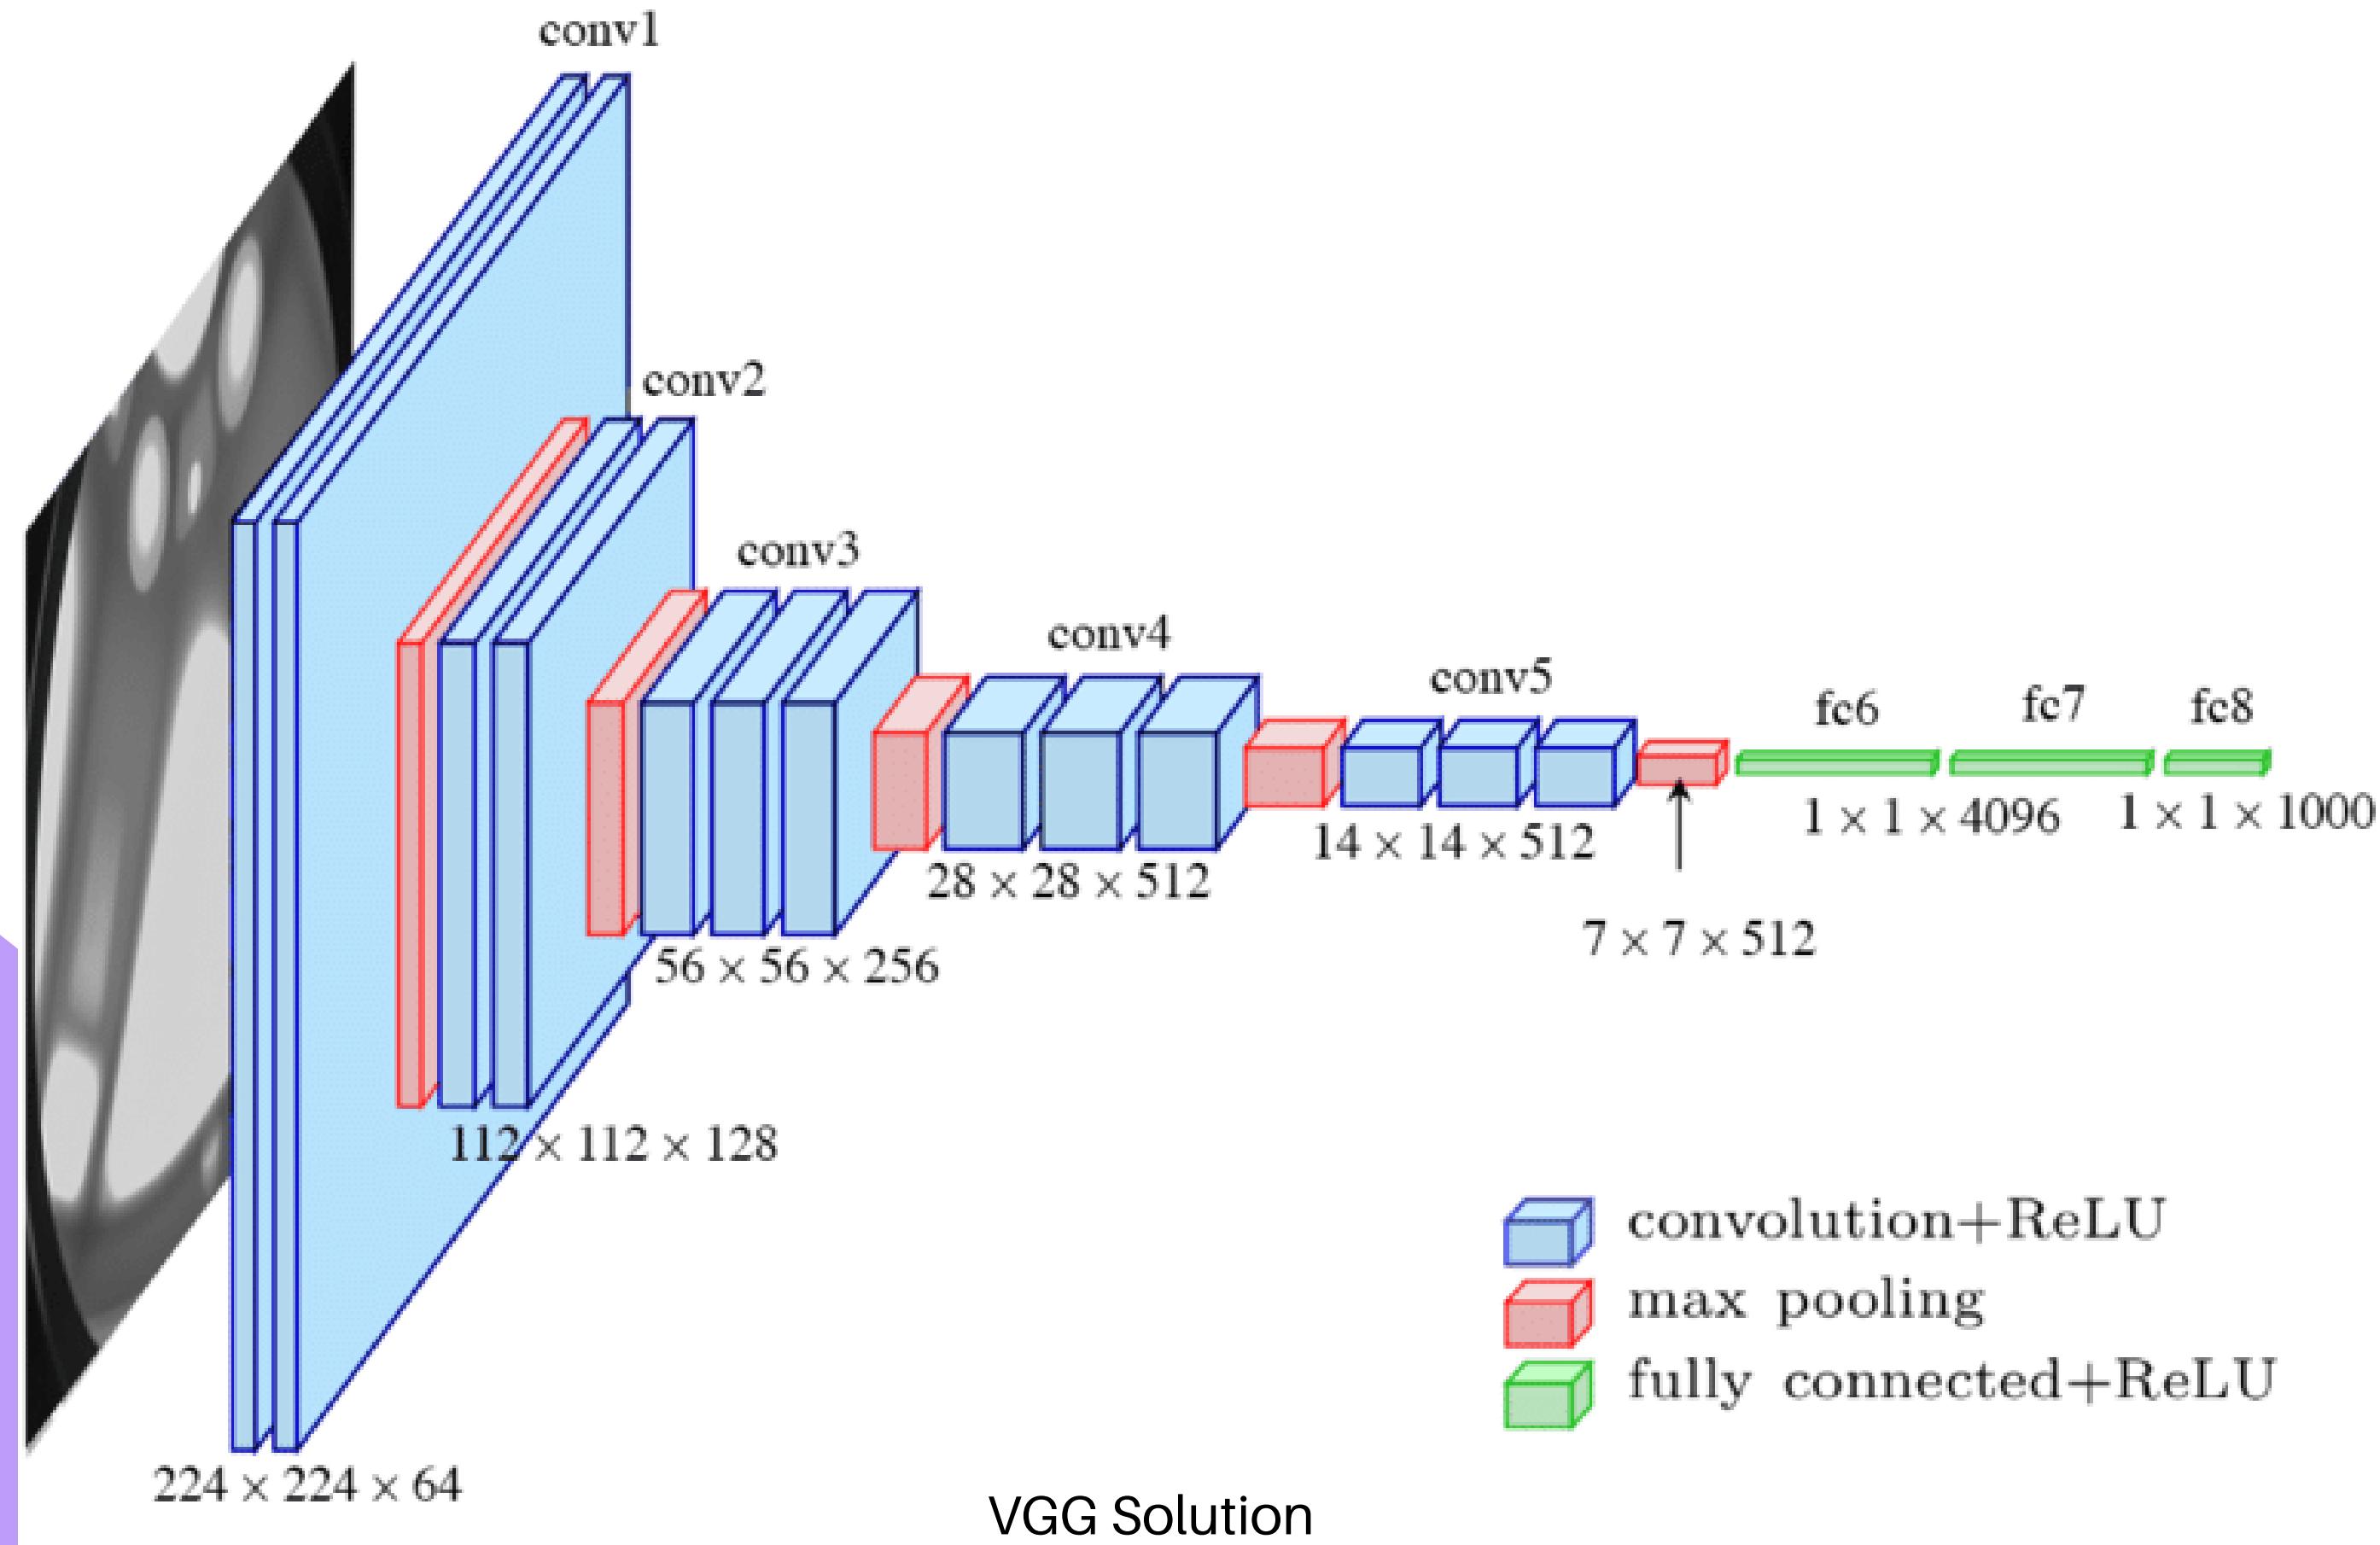


Figure 2. Model Scaling. (a) is a baseline network example; (b)-(d) are conventional scaling that only increases one dimension (either width, depth, or resolution). (e) is our proposed compound scaling method that uniformly scales all three dimensions with a fixed ratio.

EfficientNet Solution



MobileNet solution



Experiments



Experiments

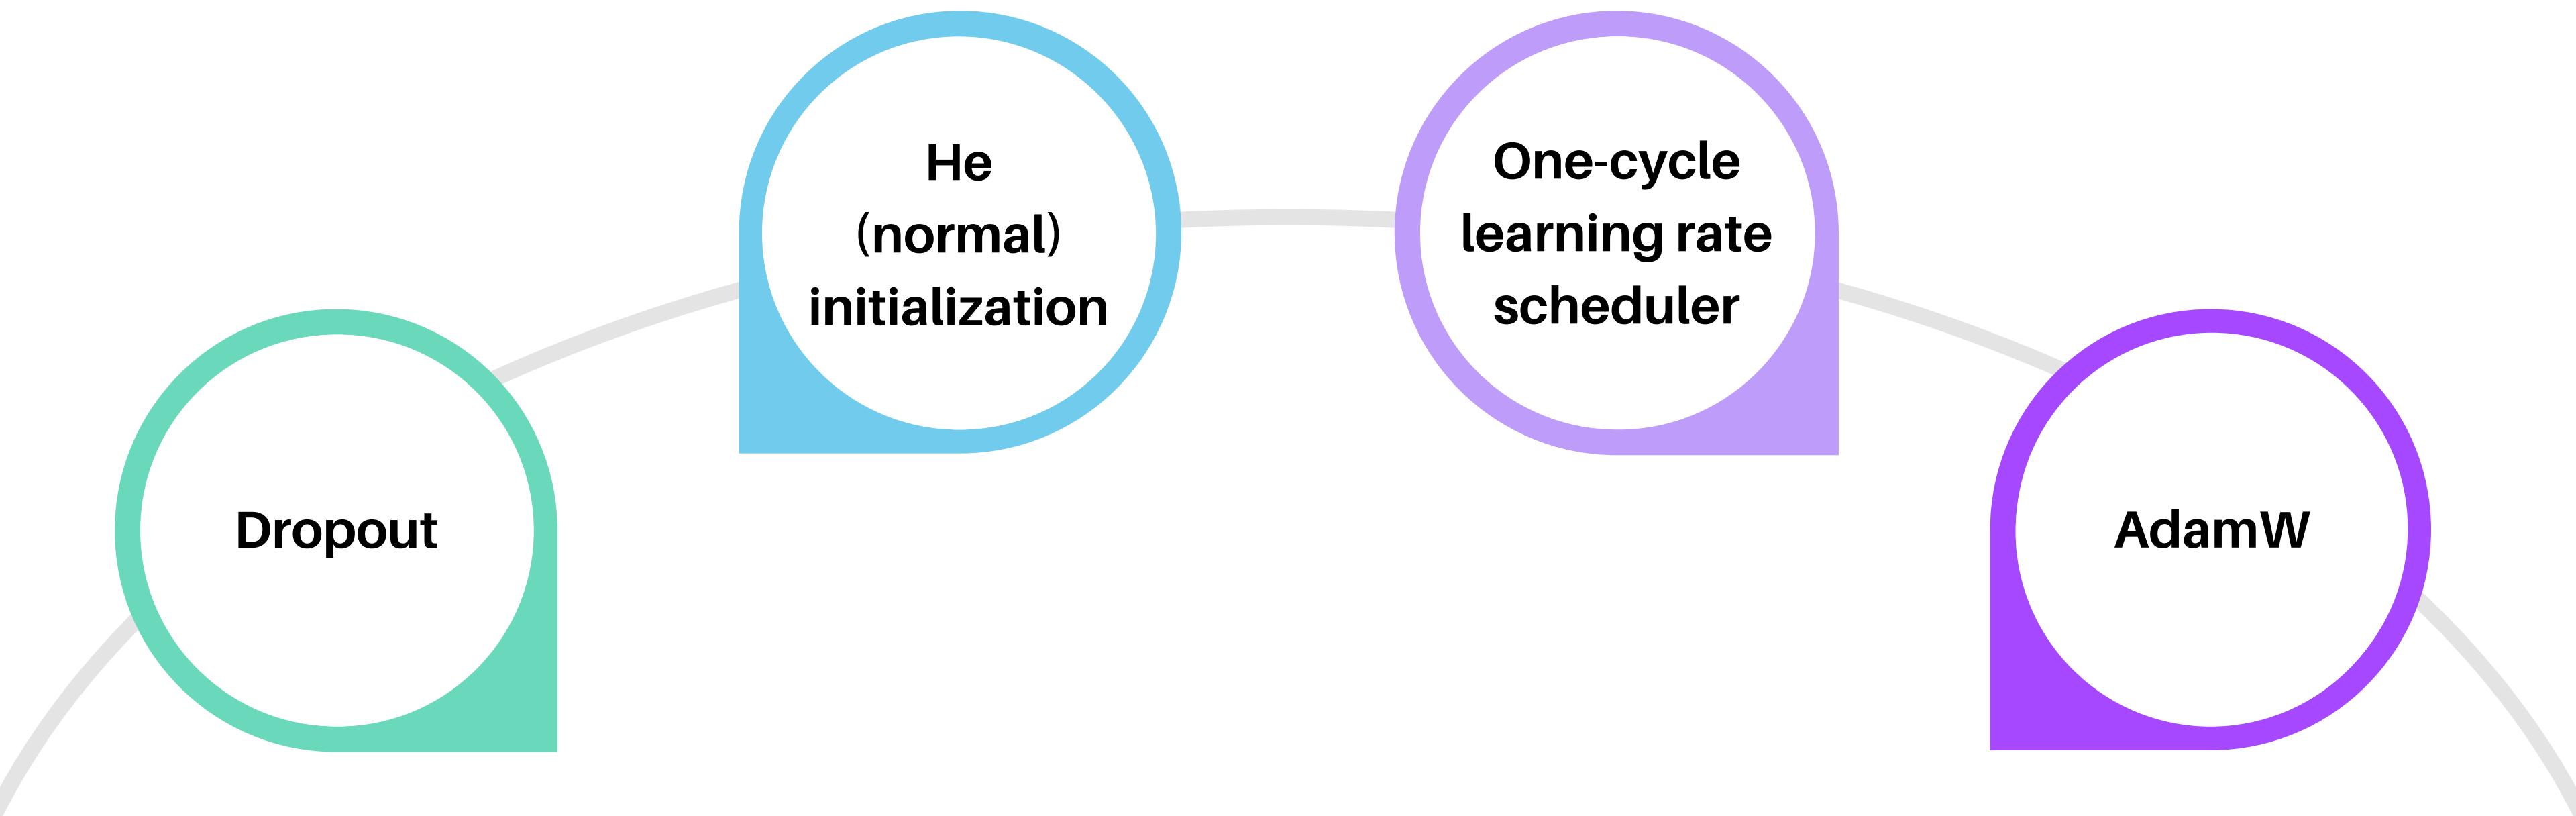
Parameters

- training epochs = 25
- batch size = 64
- patience (Early Stopping) = 8
- Adam = default values
- Cross-Entropy loss function

Evaluation metrics

- accuracy score
- confusion matrix
- precision, recall, f1-score
- visualization of wrong predictions

Optimization techniques



Baseline CNN

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656

Model	Number of parameters
Our VGG16	134,285,380
VGG16	134,276,932
Our VGG11	128,788,228
EfficientNetB7	63,797,204
ResNet50	23,516,228
Our ResNet50	23,516,228
Our ResNet18	11,178,564
DenseNet121	6,957,956
Our DenseNet121	6,957,956
Our EfficientNetB1	6,518,308
Our Baseline CNN	6,485,956
MobileNetV2	2,228,996
Our MobileNetV2	2,228,996
Baseline CNN	364,580

Model	Wrong Predictions
Our VGG11	2
Our VGG16	5
ResNet50	7
Our Baseline CNN	7
Our ResNet50	7
Baseline CNN	8
DenseNet121	8
Our DenseNet121	8
VGG16	9
Our EfficientNetB1	10
Our ResNet18	10
MobileNetV2	16
Our MobileNetV2	16
EfficientNetB7	37

DenseNet

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656

Model	Number of parameters
Our VGG16	134,285,380
VGG16	134,276,932
Our VGG11	128,788,228
EfficientNetB7	63,797,204
ResNet50	23,516,228
Our ResNet50	23,516,228
Our ResNet18	11,178,564
DenseNet121	6,957,956
Our DenseNet121	6,957,956
Our EfficientNetB1	6,518,308
Our Baseline CNN	6,485,956
MobileNetV2	2,228,996
Our MobileNetV2	2,228,996
Baseline CNN	364,580

Model	Wrong Predictions
Our VGG11	2
Our VGG16	5
ResNet50	7
Our Baseline CNN	7
Our ResNet50	7
Baseline CNN	8
DenseNet121	8
Our DenseNet121	8
VGG16	9
Our EfficientNetB1	10
Our ResNet18	10
MobileNetV2	16
Our MobileNetV2	16
EfficientNetB7	37

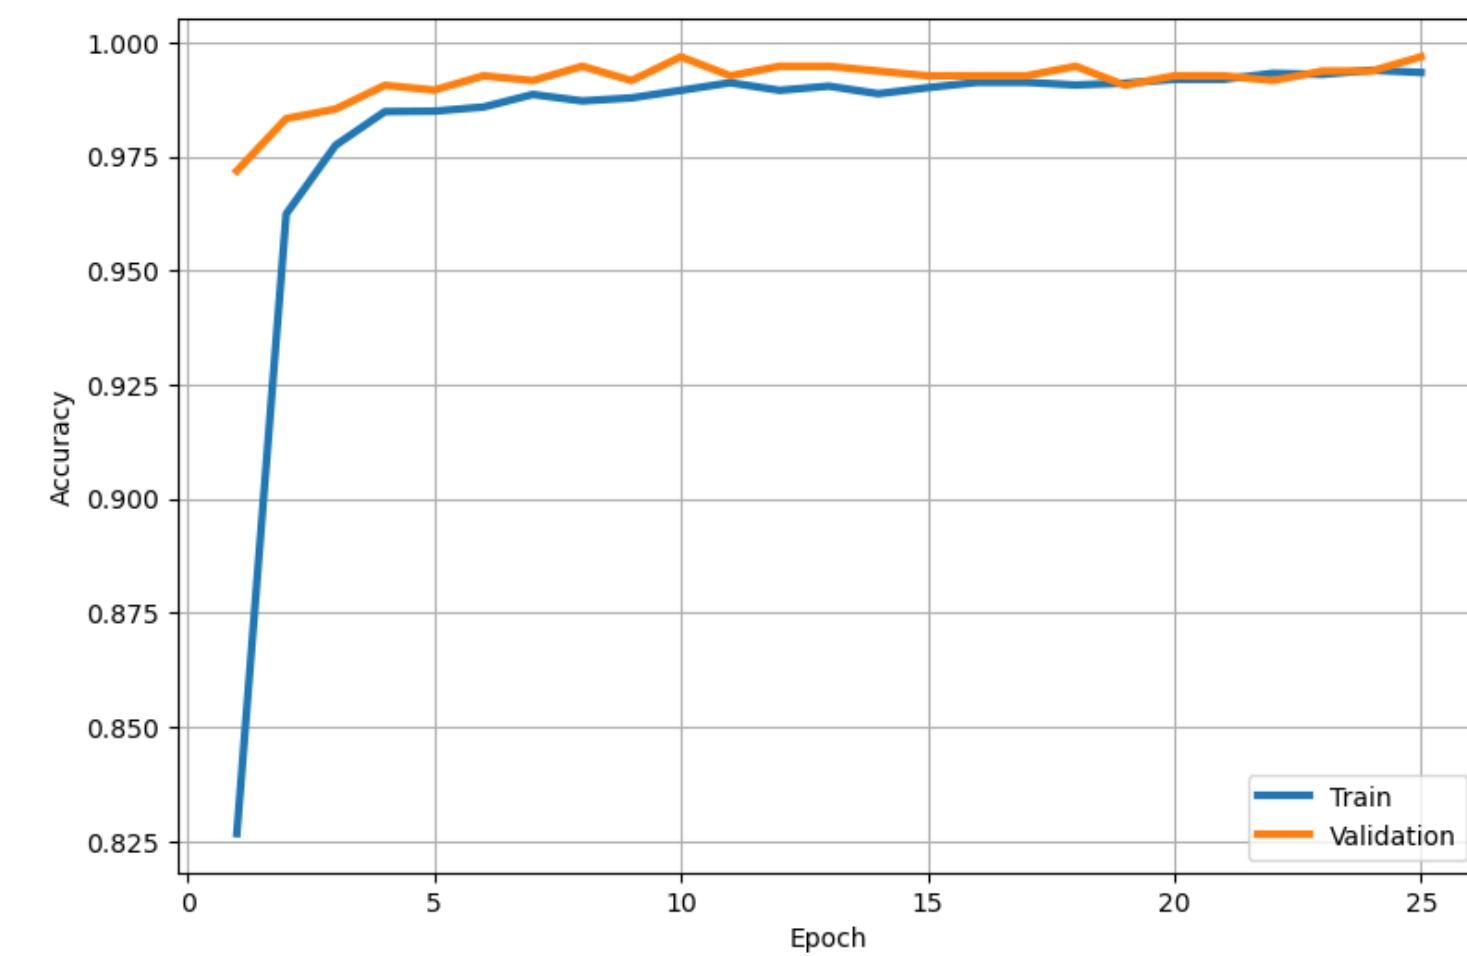
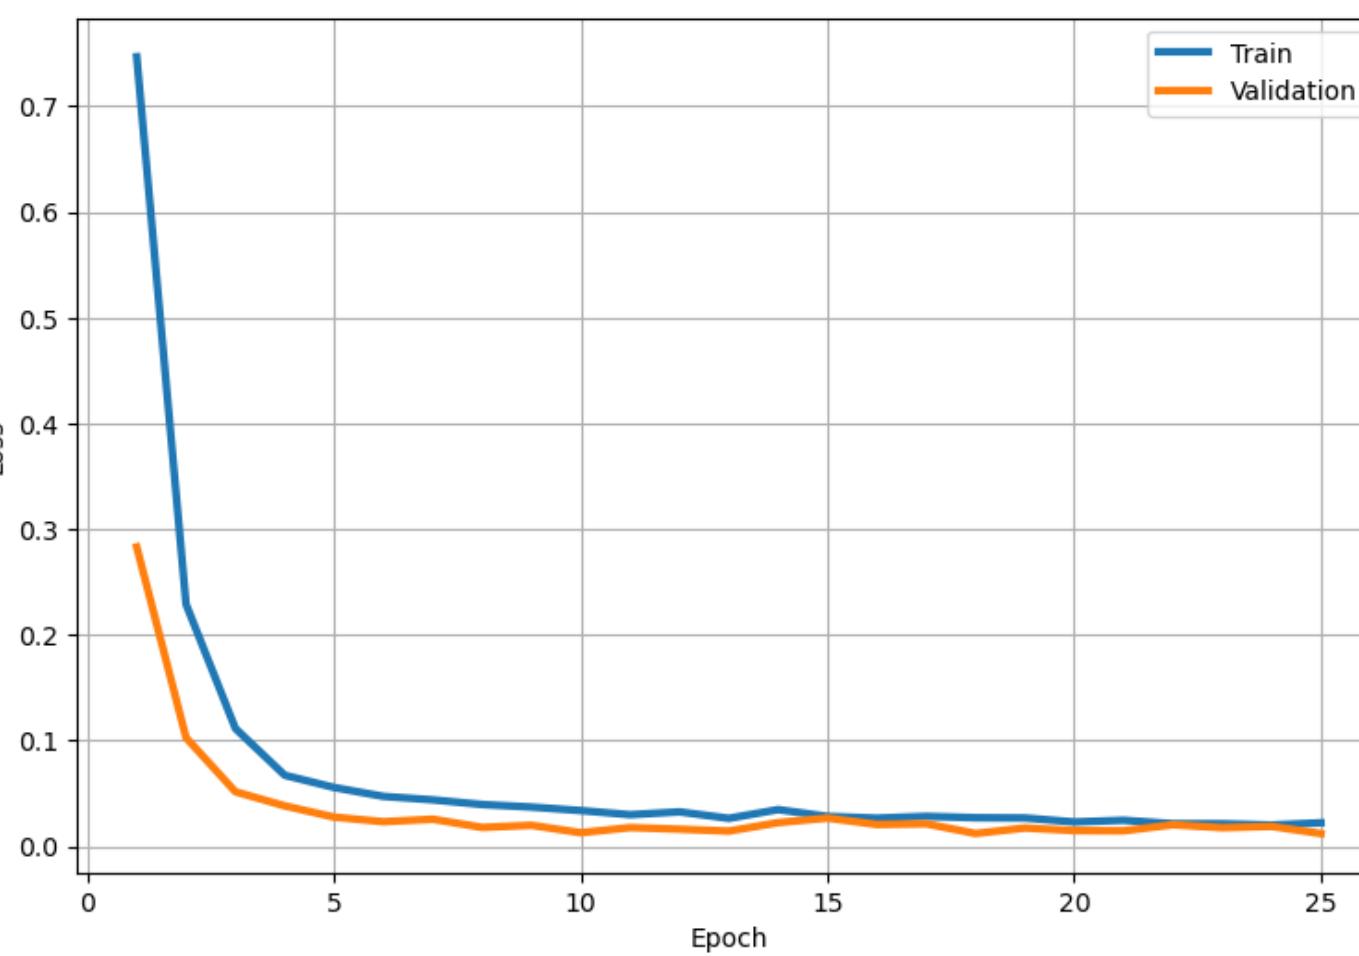
EfficientNet

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656

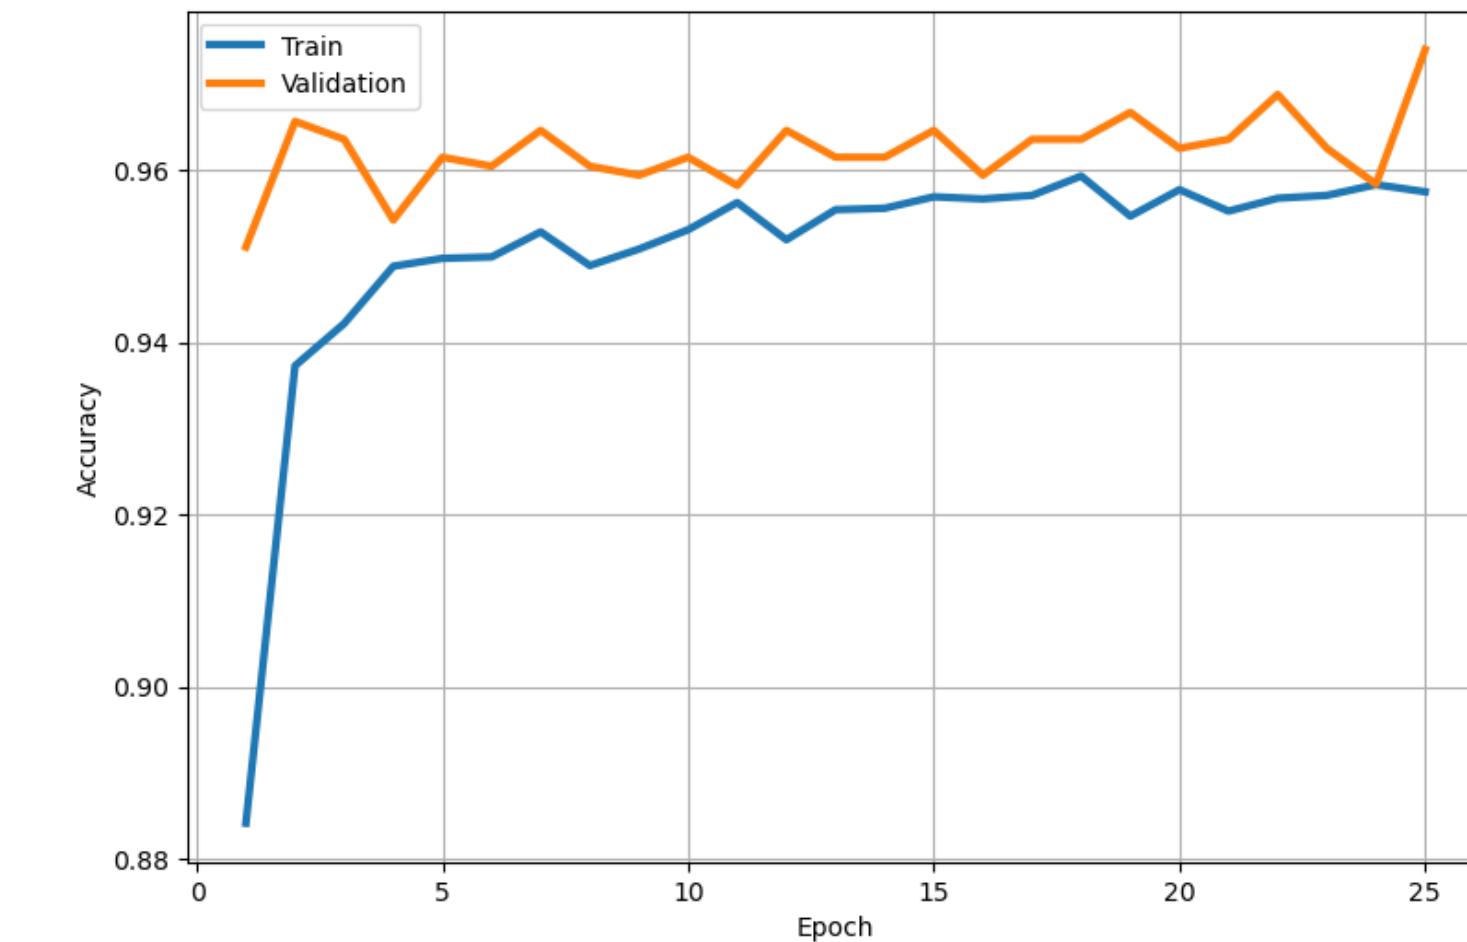
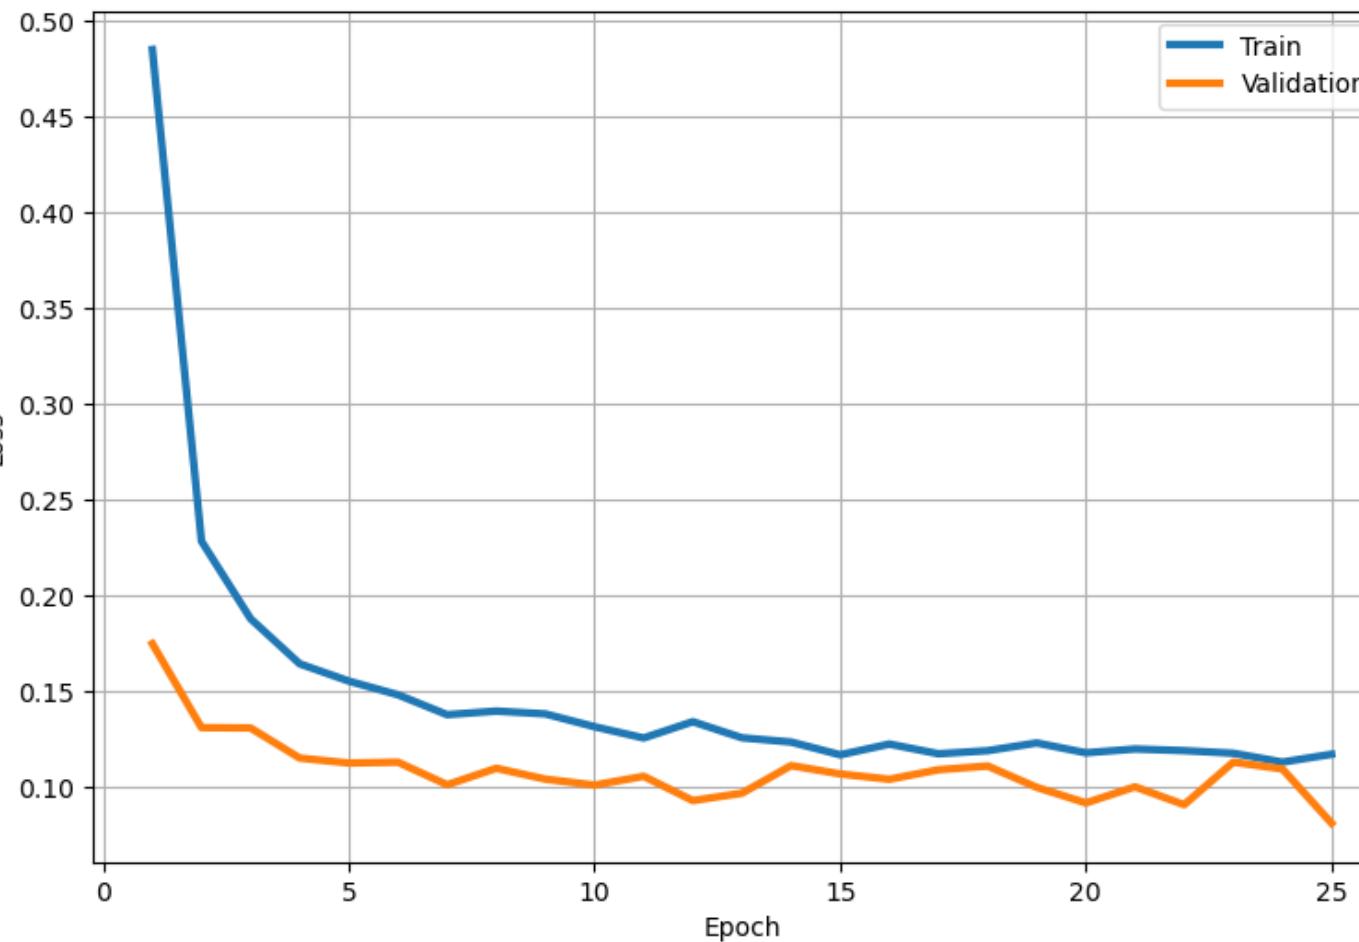
Model	Number of parameters
Our VGG16	134,285,380
VGG16	134,276,932
Our VGG11	128,788,228
EfficientNetB7	63,797,204
ResNet50	23,516,228
Our ResNet50	23,516,228
Our ResNet18	11,178,564
DenseNet121	6,957,956
Our DenseNet121	6,957,956
Our EfficientNetB1	6,518,308
Our Baseline CNN	6,485,956
MobileNetV2	2,228,996
Our MobileNetV2	2,228,996
Baseline CNN	364,580

Model	Wrong Predictions
Our VGG11	2
Our VGG16	5
ResNet50	7
Our Baseline CNN	7
Our ResNet50	7
Baseline CNN	8
DenseNet121	8
Our DenseNet121	8
VGG16	9
Our EfficientNetB1	10
Our ResNet18	10
MobileNetV2	16
Our MobileNetV2	16
EfficientNetB7	37

Our (EfficientNetB1)



Original (EfficientNetB7)



MobileNet

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656

Model	Number of parameters
Our VGG16	134,285,380
VGG16	134,276,932
Our VGG11	128,788,228
EfficientNetB7	63,797,204
ResNet50	23,516,228
Our ResNet50	23,516,228
Our ResNet18	11,178,564
DenseNet121	6,957,956
Our DenseNet121	6,957,956
Our EfficientNetB1	6,518,308
Our Baseline CNN	6,485,956
MobileNetV2	2,228,996
Our MobileNetV2	2,228,996
Baseline CNN	364,580

Model	Wrong Predictions
Our VGG11	2
Our VGG16	5
ResNet50	7
Our Baseline CNN	7
Our ResNet50	7
Baseline CNN	8
DenseNet121	8
Our DenseNet121	8
VGG16	9
Our EfficientNetB1	10
Our ResNet18	10
MobileNetV2	16
Our MobileNetV2	16
EfficientNetB7	37

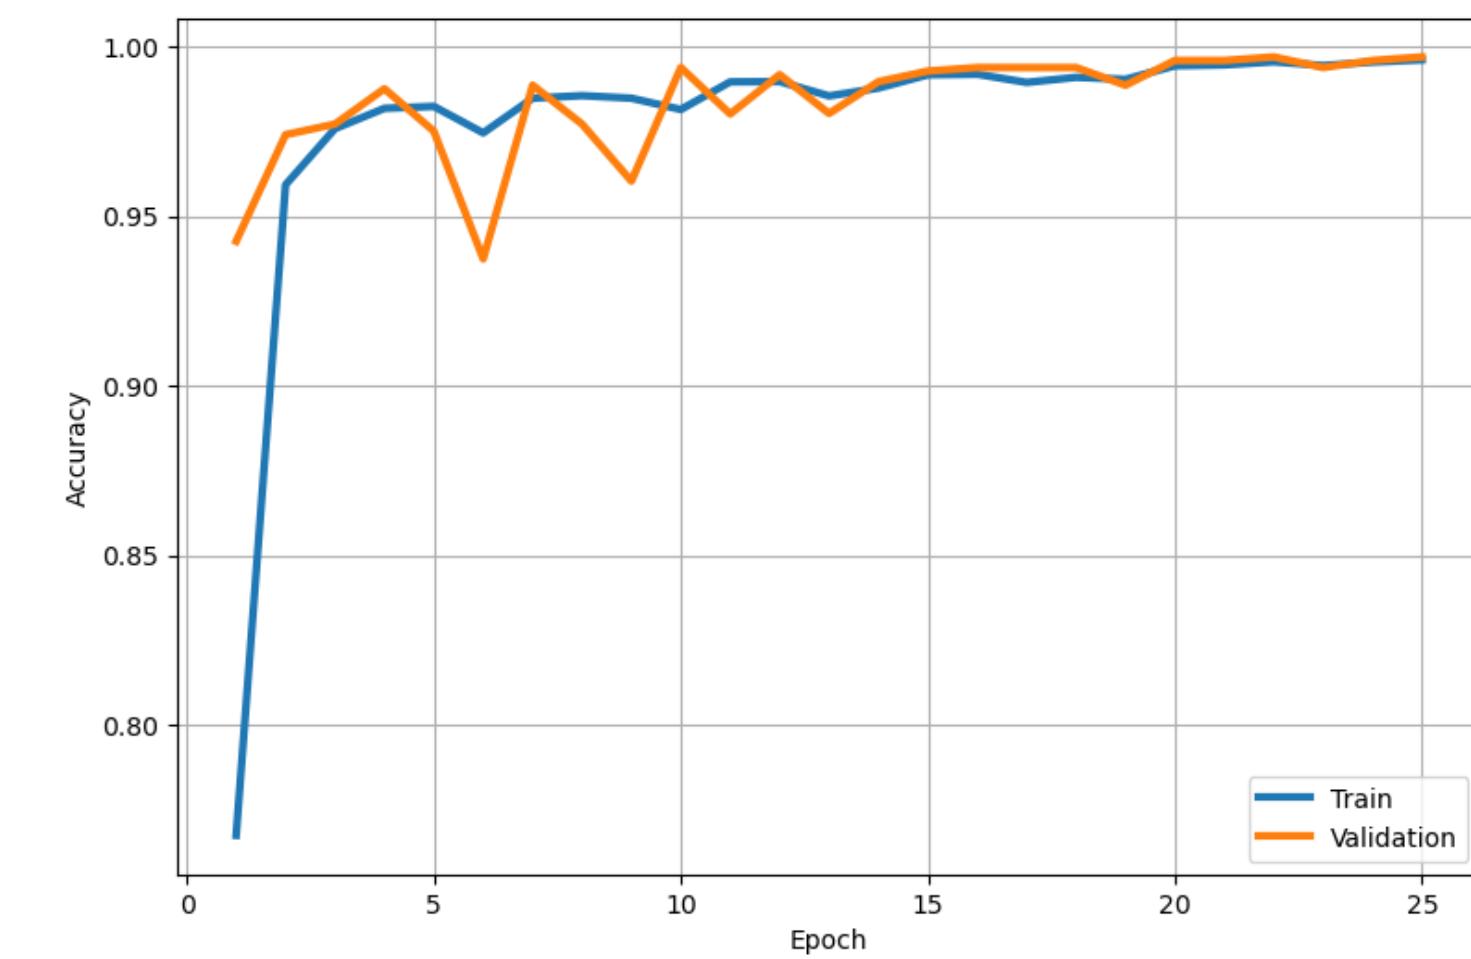
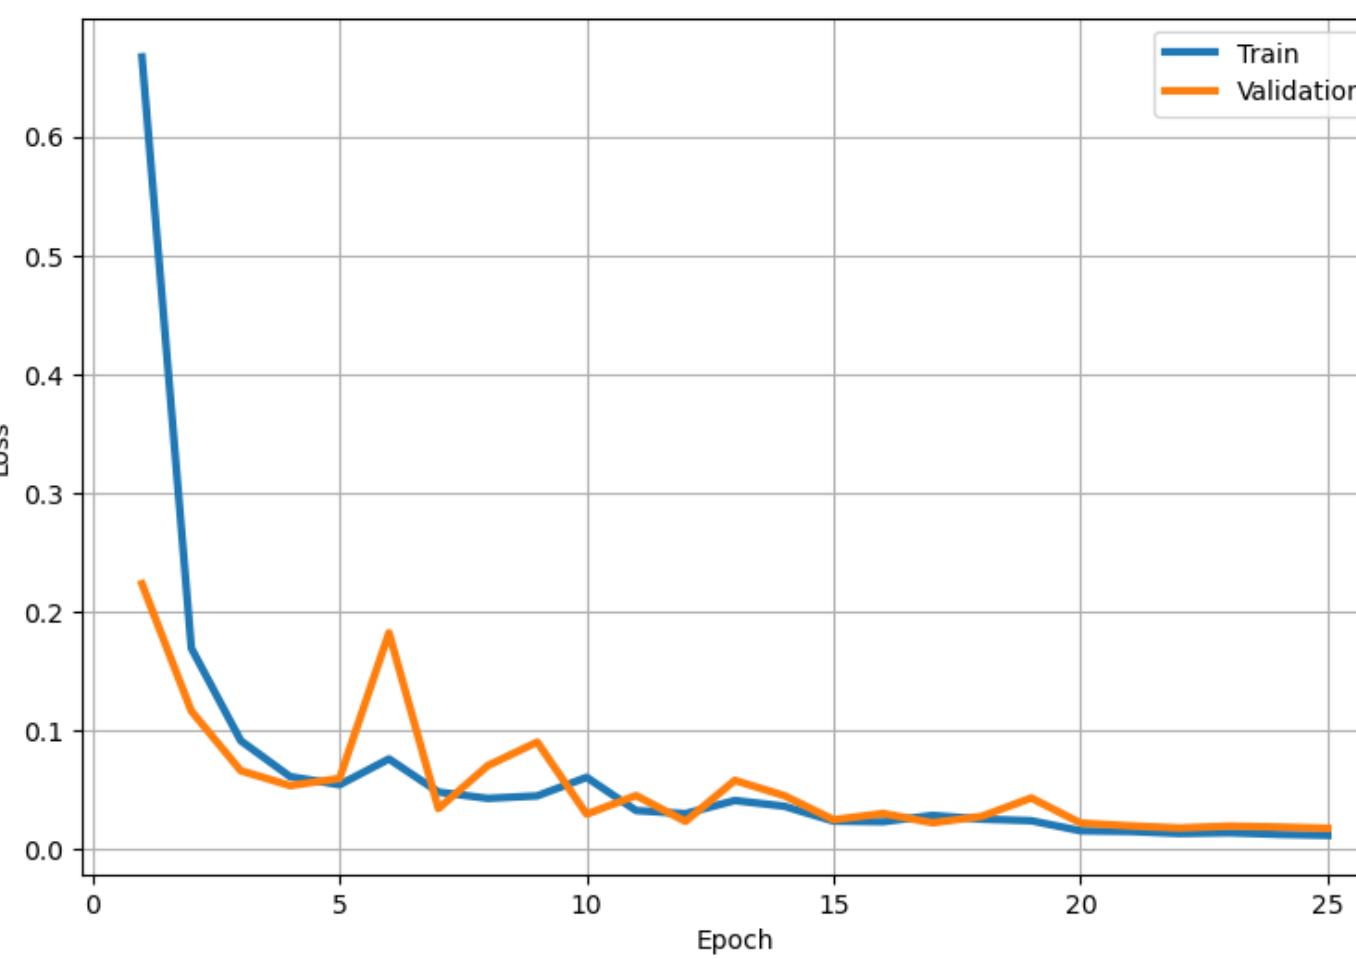
ResNet

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656

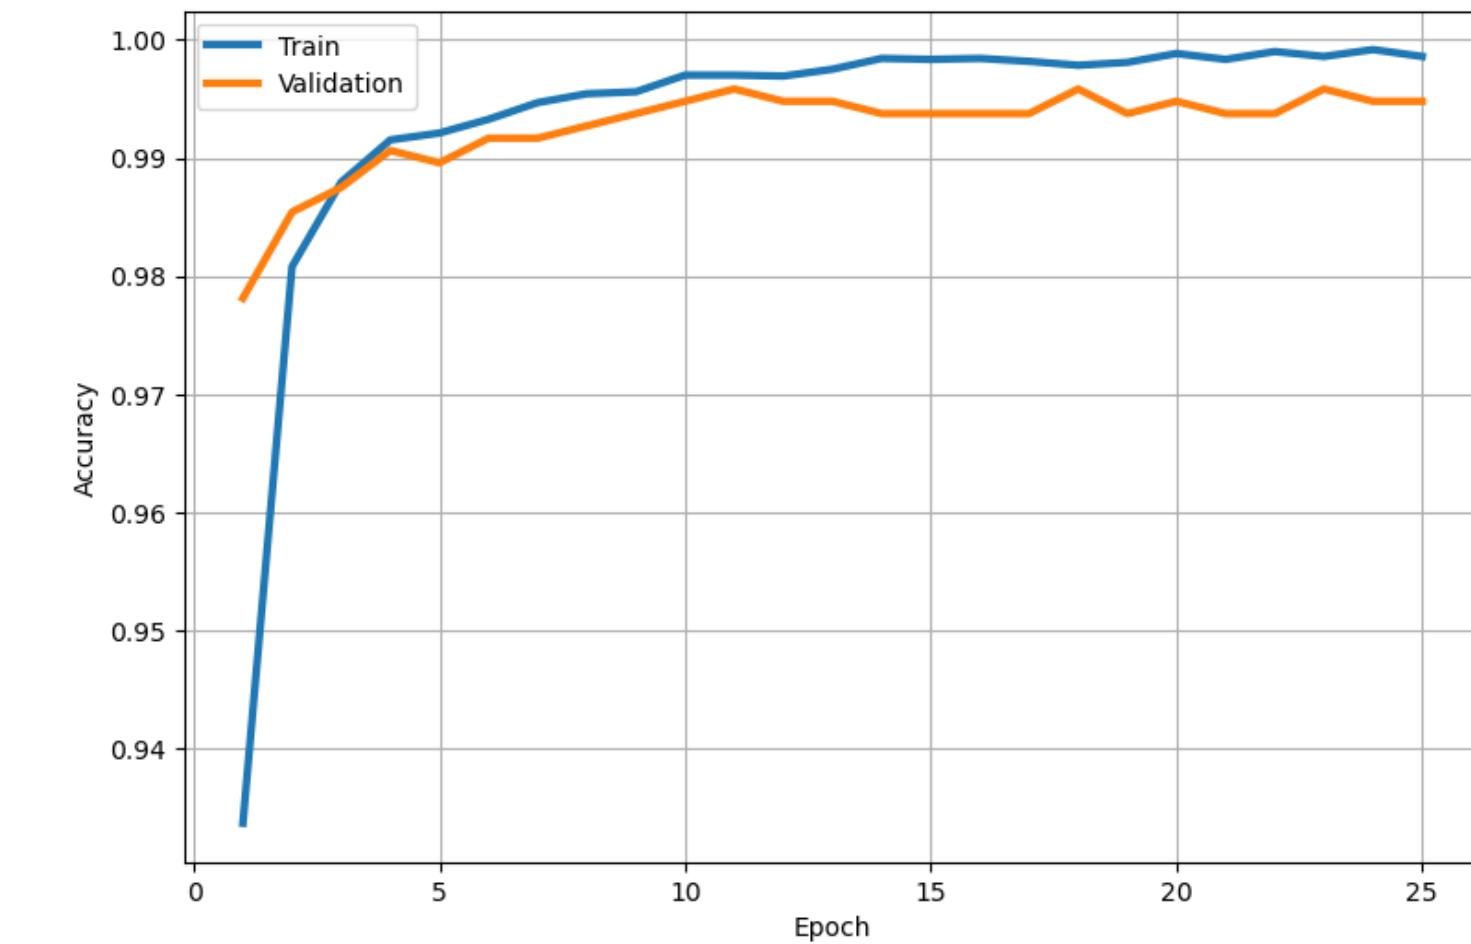
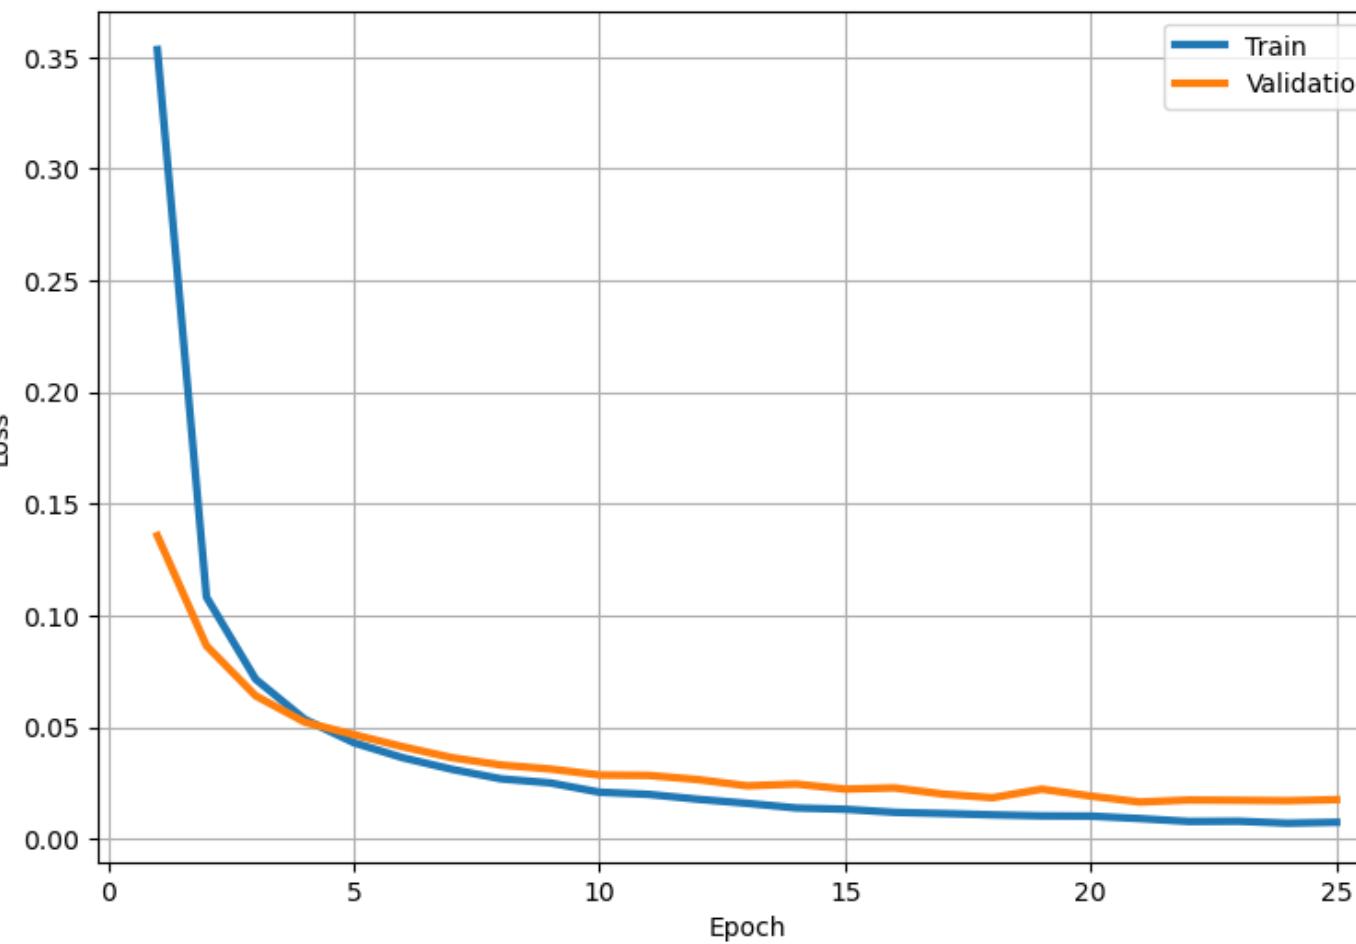
Model	Number of parameters
Our VGG16	134,285,380
VGG16	134,276,932
Our VGG11	128,788,228
EfficientNetB7	63,797,204
ResNet50	23,516,228
Our ResNet50	23,516,228
Our ResNet18	11,178,564
DenseNet121	6,957,956
Our DenseNet121	6,957,956
Our EfficientNetB1	6,518,308
Our Baseline CNN	6,485,956
MobileNetV2	2,228,996
Our MobileNetV2	2,228,996
Baseline CNN	364,580

Model	Wrong Predictions
Our VGG11	2
Our VGG16	5
ResNet50	7
Our Baseline CNN	7
Our ResNet50	7
Baseline CNN	8
DenseNet121	8
Our DenseNet121	8
VGG16	9
Our EfficientNetB1	10
Our ResNet18	10
MobileNetV2	16
Our MobileNetV2	16
EfficientNetB7	37

Our (ResNet18)



Original (ResNet50)



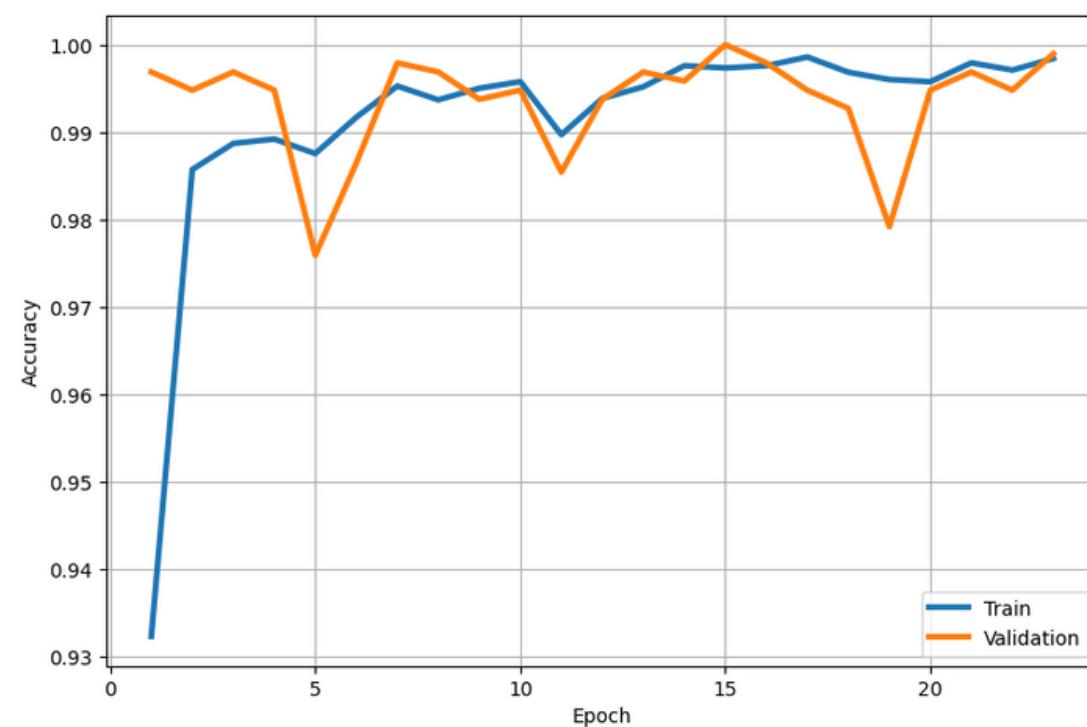
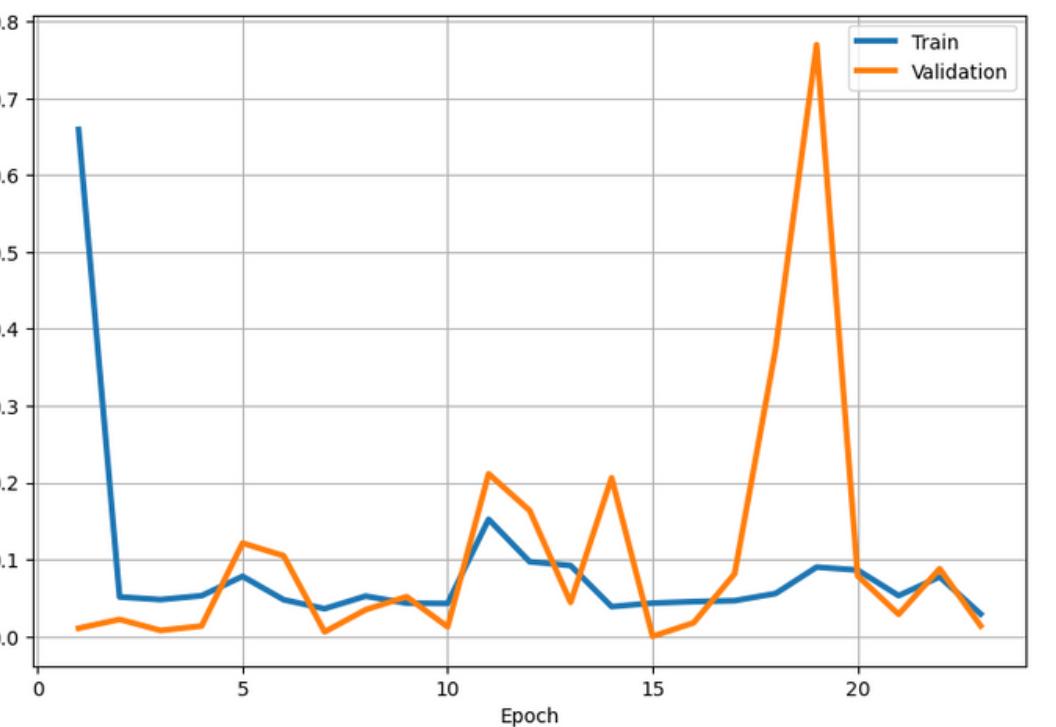
VGG

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656

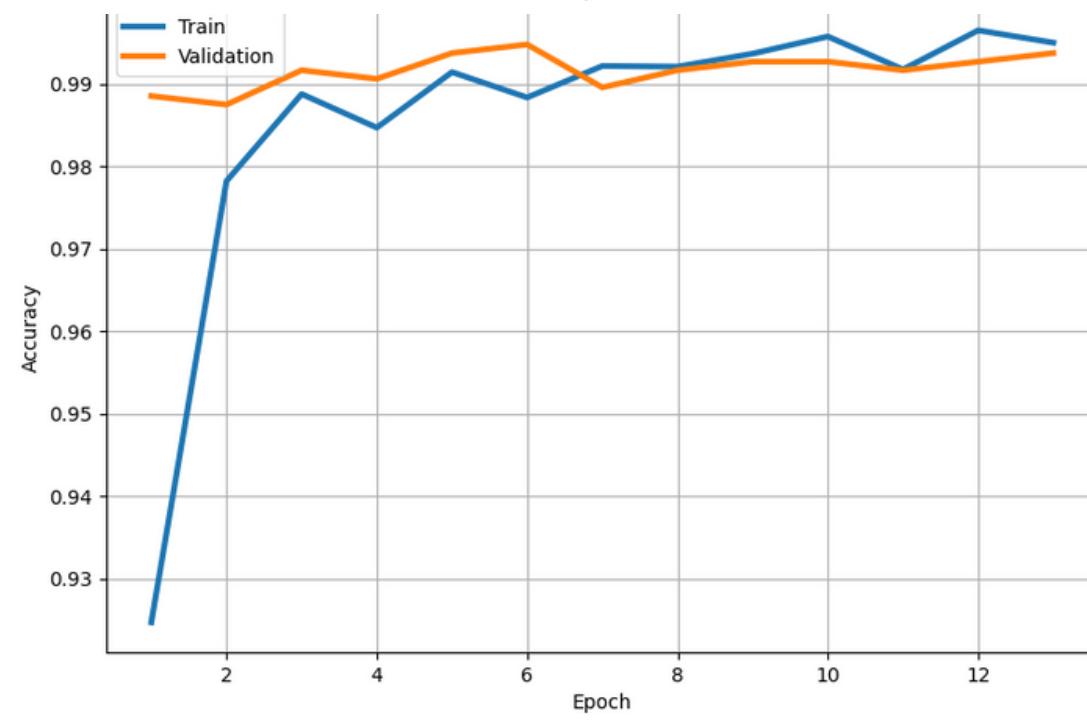
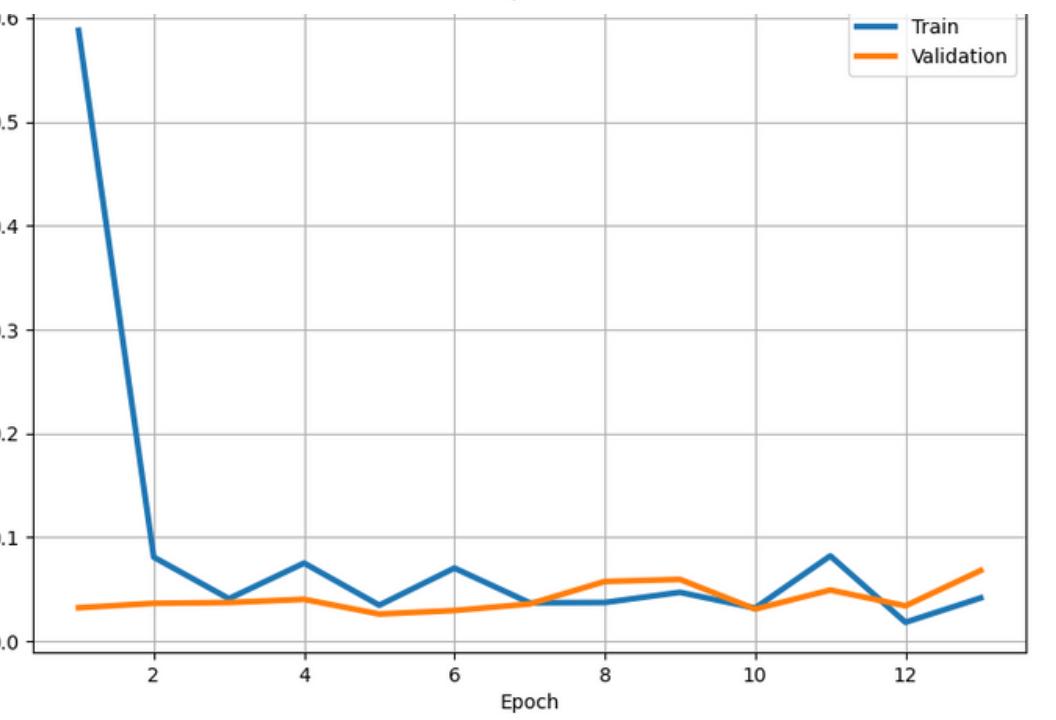
Model	Number of parameters
Our VGG16	134,285,380
VGG16	134,276,932
Our VGG11	128,788,228
EfficientNetB7	63,797,204
ResNet50	23,516,228
Our ResNet50	23,516,228
Our ResNet18	11,178,564
DenseNet121	6,957,956
Our DenseNet121	6,957,956
Our EfficientNetB1	6,518,308
Our Baseline CNN	6,485,956
MobileNetV2	2,228,996
Our MobileNetV2	2,228,996
Baseline CNN	364,580

Model	Wrong Predictions
Our VGG11	2
Our VGG16	5
ResNet50	7
Our Baseline CNN	7
Our ResNet50	7
Baseline CNN	8
DenseNet121	8
Our DenseNet121	8
VGG16	9
Our EfficientNetB1	10
Our ResNet18	10
MobileNetV2	16
Our MobileNetV2	16
EfficientNetB7	37

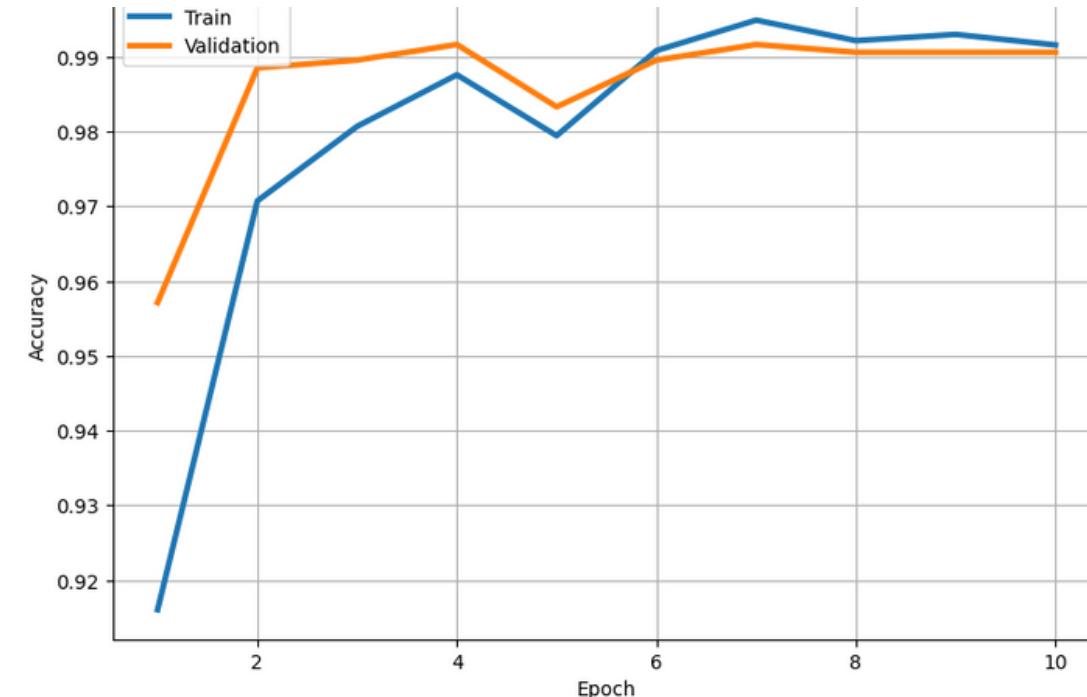
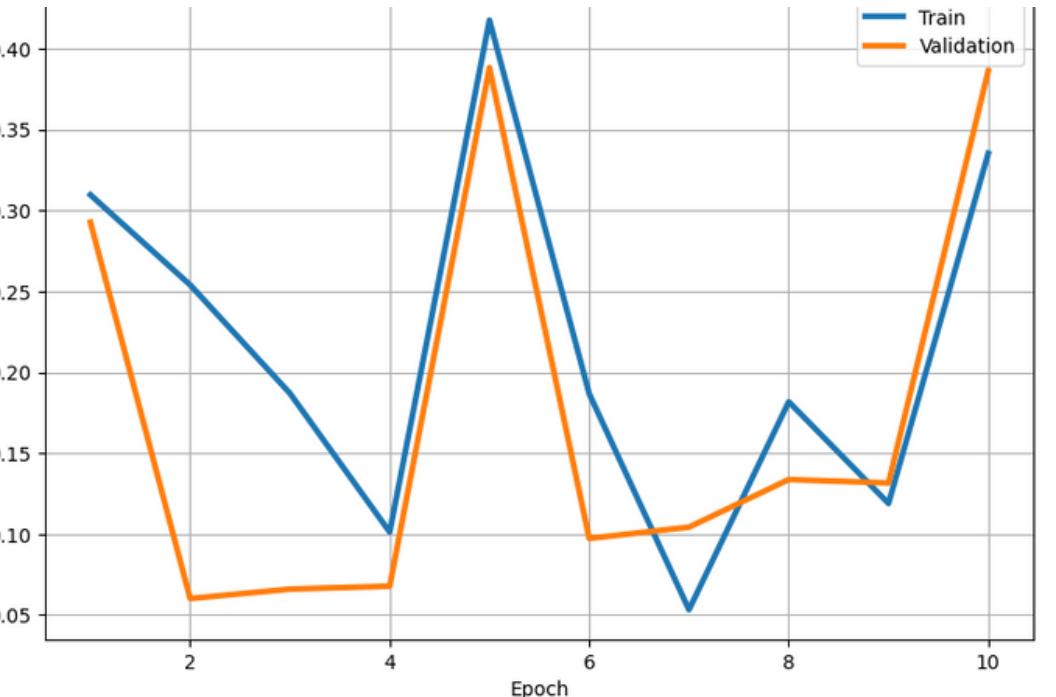
Our
(VGG11)



Our
(VGG16BN)

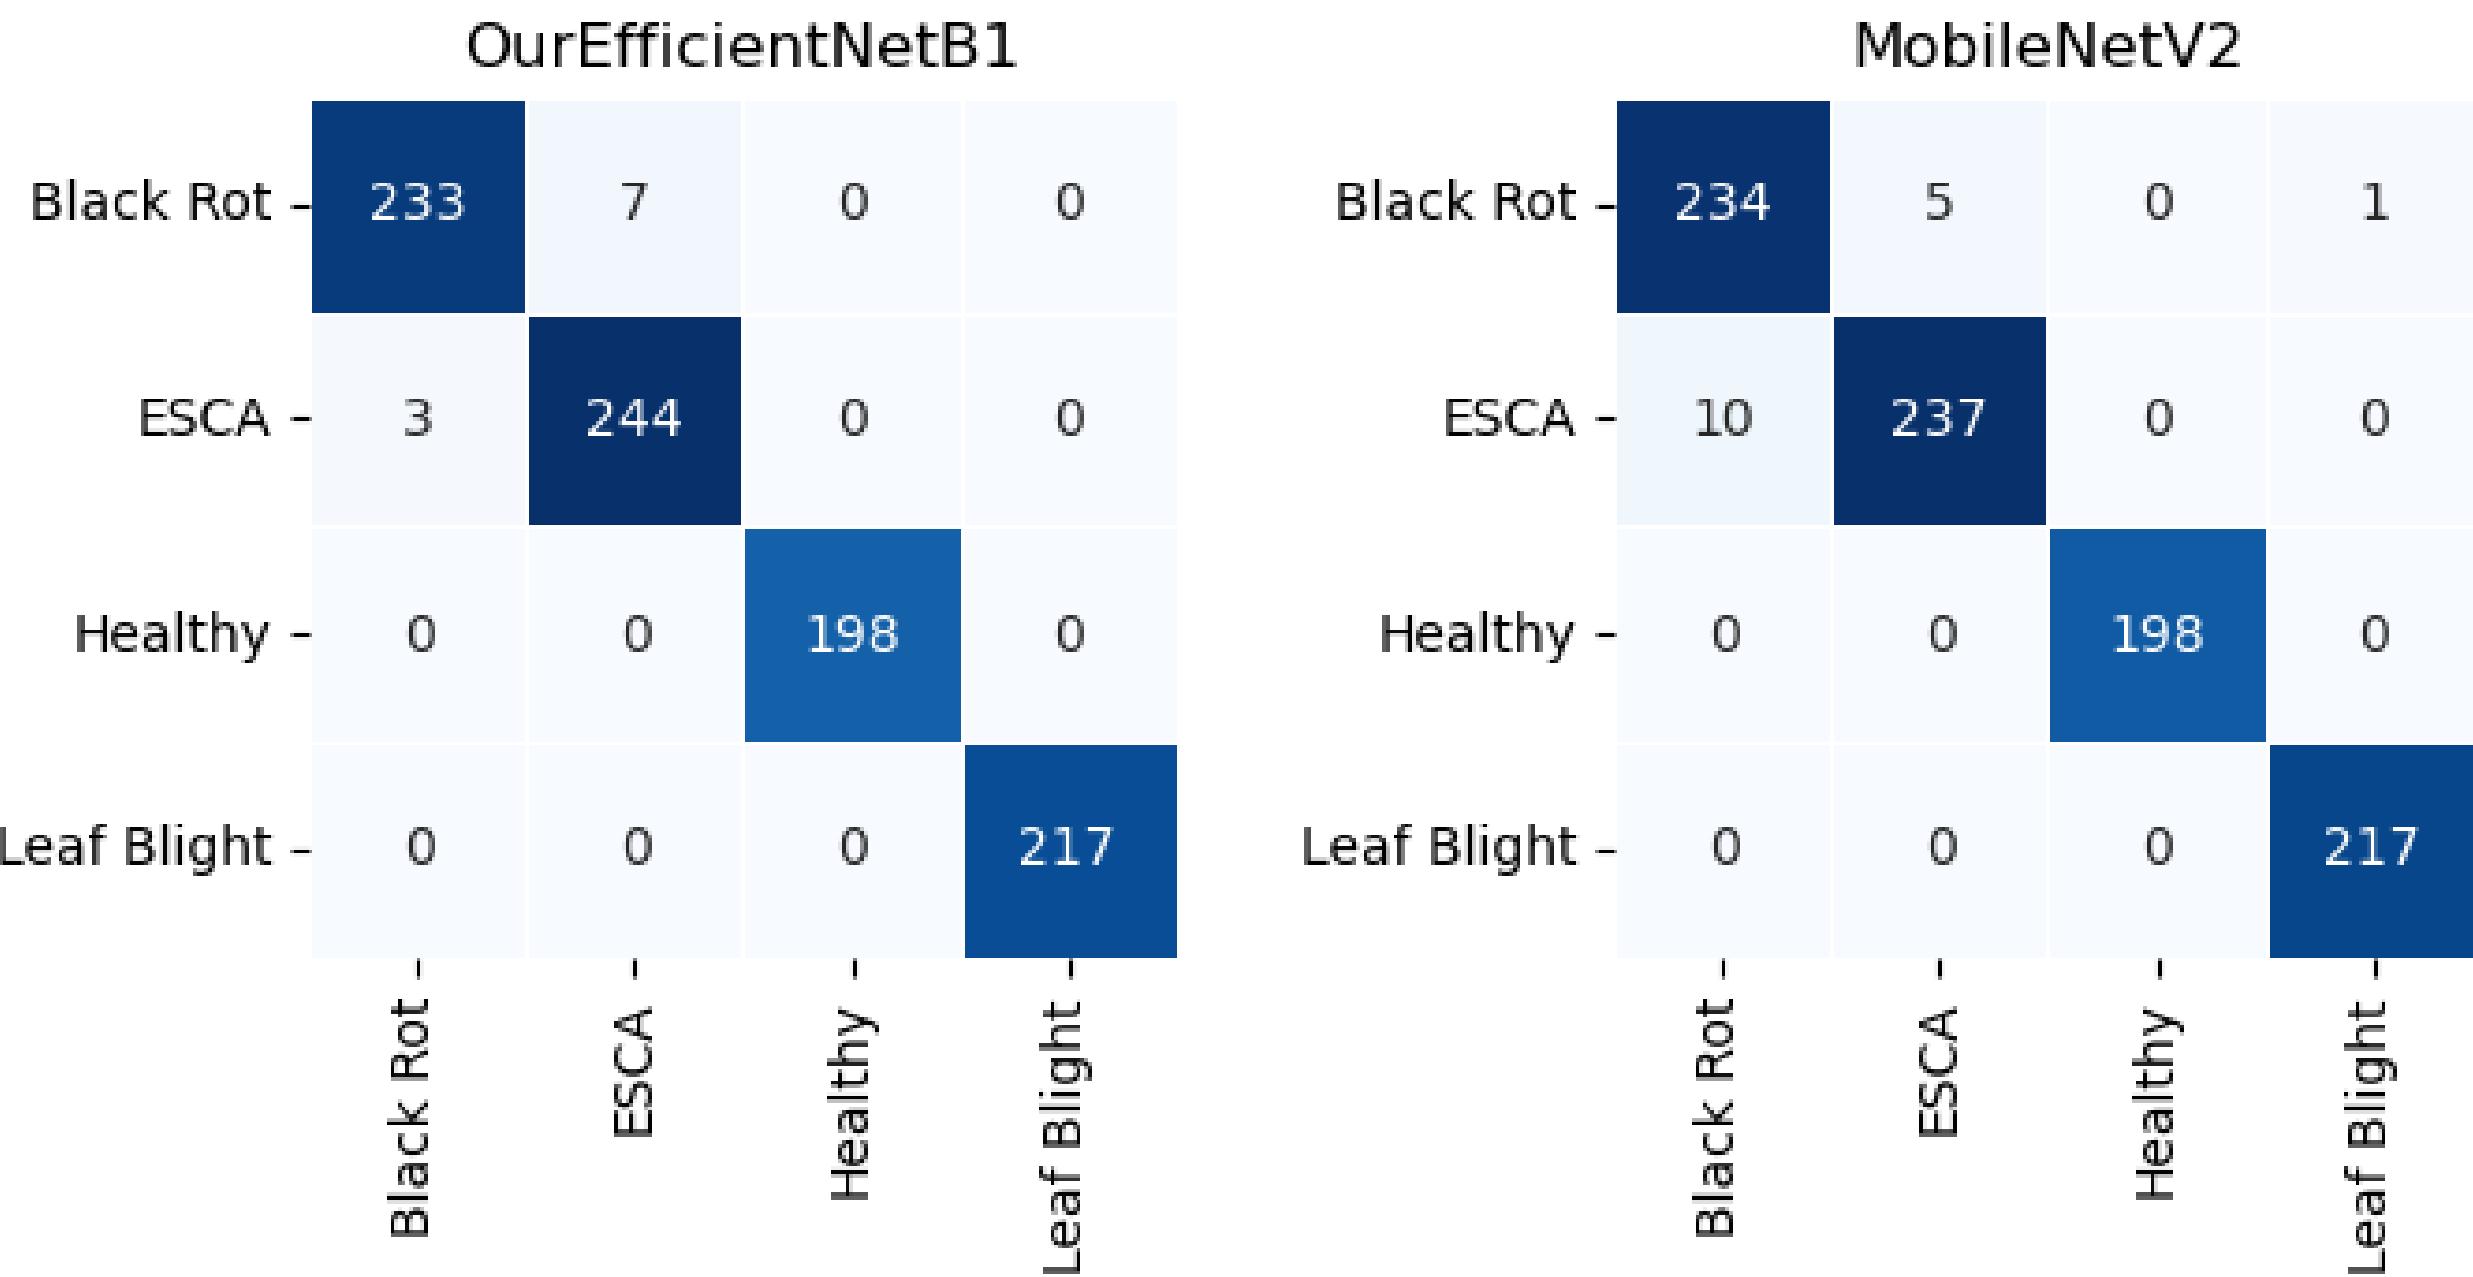


Original
(VGG16)



Summing up...

Model	Accuracy
Our VGG11	0.9979
Our VGG16	0.9948
ResNet50	0.9927
Our Baseline CNN	0.9927
Our EfficientNetB1	0.9927
Our ResNet50	0.9927
Baseline CNN	0.9917
VGG16	0.9906
Our ResNet18	0.9896
MobileNetV2	0.9833
DenseNet121	0.9816
Our MobileNetV2	0.9812
Our DenseNet121	0.9795
EfficientNetB7	0.9656



Max-voting ensemble

(using the best 3 models)

Our

- VGG11
- Baseline CNN
- EfficientNetB1
- Accuracy score: 99.56%
- Misclassifications: 4

Original

- VGG16
- Baseline CNN
- ResNet50
- Accuracy score: 98.89%
- Misclassifications: 10

Our ensemble

Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Leaf Blight



Predicted: Healthy
Actual: Black Rot



Original ensemble

Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Black Rot



Predicted: Black Rot
Actual: ESCA



Predicted: Black Rot
Actual: ESCA



Predicted: ESCA
Actual: Black Rot



Predicted: ESCA
Actual: Black Rot



Conclusion

What we have done?

- analyzed, implemented, and improved a specific *Kaggle* project
- evaluated many pre-trained models for classifying 4 grape diseases
- translated the code from *TensorFlow* to *PyTorch*
- adjusted the original dataset splitting, experimented with new models, and applied novel optimization techniques

Future directions

- need more comprehensive and robust models
- create a new standard and large dataset (data availability acts as a bottleneck)